

Advanced Algorithms Notes

Alberto Lazari

II Semester – 2023

Index

1 – DFS (Depth First Search)	3
1.1 – Applications	3
2 – BFS (Breadth First Search)	3
3 – MST (Minimum[-weight] Spanning Tree)	3
3.1 – Prim	3
3.1.1 – Prim with heap	3
3.2 – Kruskal	4
3.2.1 – Kruskal with disjoint sets	4
4 – SS (Single-Source) Shortest Paths	4
4.1 – Non-negative weights – Dijkstra	4
4.2 – General case – Bellman-Ford	4
5 – AP (All Pair) Shortest Paths	5
5.1 – Floyd-Warshall	5
6 – Maximum flows	5
6.1 – Definitions	5
6.2 – Ford-Fulkerson	5
7 – NP-hardness	5
7.1 – Reduction	6
7.2 – NP-hard Problems	6
8 – Approximation algorithms	6
8.1 – Vertex cover	6
8.2 – Metric TSP	7
8.2.1 – Eulerian circuit approach	7
8.3 – Set cover	7
9 – Randomized algorithms	7

10 – Minimum cut – Karger	8
10.1 – Algorithm	8
10.2 – Definitions	8
10.2.1 – Multigraphs	8
10.2.2 – Full Contraction	8
10.3 – Karger-Stein	8
10.4 – 2020 version	8

Lecture 2

1 – DFS (Depth First Search)

Complexity $O(n + m)$

1.1 – Applications

Derived using DFS (or BFS) in $O(n + m)$

- Path between source vertex s to arbitrary t : add a parent field to vertices.
When t is found return the parents backtrace
- Find cycle: use parent field on vertices and ancestor on edges
- Connected components:
 1. run DFS (or BFS) n times
 2. Keep a counter k to increment on every “untouched” source vertex
 3. Assign k to v .id, instead of 1 \rightarrow label vertexes based on its component
 4. If at the end $k > 1$, then multiple components were found

Lecture 3

2 – BFS (Breadth First Search)

Complexity $O(n + m)$

3 – MST (Minimum[-weight] Spanning Tree)

MST $(G = (V, E), s)$

Tree created from a source vertex s , the root of the tree

Lecture 4

3.1 – Prim

Complexity $O(m \cdot n)$

Make cuts to separate a growing set A (initialized to $\{s\}$), and find *light edges*. Add the light edge found with the cut to A and repeat, until you have a tree (no more vertices outside $V \setminus A$)

The search for the light edge is $O(m)$ and is repeated n times, but it can be optimized

3.1.1 – Prim with heap

Complexity $O(m \log n)$

Use a heap to store vertices, ordered on their cost to reach from a vertex already processed (light edge that crosses the cut) For every vertex that you

put in A (actually that you extract from the heap H) check if you can update the cost of the vertices still in H

In order to keep trace of the actual edges, instead of the vertices, it's needed to save the parent of every vertex you update

The complexity is actually $O(n \log n + m \log n)$, but graph G is connected $\Rightarrow m \geq n - 1$

Lecture 5

3.2 – Kruskal

Complexity $O(m \cdot n)$ (when implemented with adjacency list, because of frequent cycle checks)

Extremely simple:

1. A is an empty forest;
2. Sort E by weight (ascending order);
3. If adding $e \in E$ to A keeps it a forest (doesn't introduce cycles) add it

3.2.1 – Kruskal with disjoint sets

Complexity $O(m \log n)$ (same of Prim with heap)

Use union-find data structure: connected components are disjoint sets to join in $O(\log n)$ time. Finds if a node is in a set in $O(\log n)$ time \Rightarrow cycle checks in logarithmic time

It's still an open problem to find MST implementation in $O(m)$

Lecture 7

4 – SS (Single-Source) Shortest Paths

SSSP ($G = (V, E), s \in V$), where G directed, weighted graph

Returns: $\text{len}(v) = \text{dist}(s, v), \forall v \in V$

4.1 – Non-negative weights – Dijkstra

Complexity $O(m \cdot n)$

Complexity can be lowered to $O((m + n) \log n)$ with heaps, similar to Prim

Lecture 8

4.2 – General case – Bellman-Ford

Complexity $O(m \cdot n)$

Need to forbid negative cycles in shortest paths, they lead to infinitely small paths → doesn't even make sense to speak about shortest paths

Bellman-Ford returns either SSSP (G, s) or a declaration that G has a negative cycle

Refine every shortest path every iteration (check every edge). In $n - 1$ iterations it reaches a fix-point. If it doesn't it means a negative cycle exist

In 2022 a **near-linear** algorithm was found

5 – AP (All Pair) Shortest Paths

Returns: $\text{dist}(v, u), \forall v, u \in V$

Running Bellman-Ford n times have complexity $O(m \cdot n^2)$. With dynamic programming complexity can be reduced up to $O(n^3 \log n)$

5.1 – Floyd-Warshall

Complexity $O(n^3)$

Iterate on 3 vertices $u, v, k \in V$ in 3 nested loops, testing whether using k in the path is better

To catch negative cycles it's sufficient to check that $\text{dist}(v, v) \geq 0, \forall v \in V$

Lecture 10

6 – Maximum flows

6.1 – Definitions

Flow network graph where edges have a capacity $c : E \rightarrow \mathbb{R}^+$.

A source s and a sink t are specified

Flow $f : E \rightarrow \mathbb{R}^+, |f| = \sum_{(s,v) \in E} f(s, v)$, basically the flow on the first edges

Flow is conserved through the graph and has to be \leq than capacity for all edges

6.2 – Ford-Fulkerson

Complexity $O(m \cdot |f^*|)$, where $|f^*|$: maximum flow

Lecture 11

7 – NP-hardness

Similar polynomial and NP-hard problems:

- Eulerian vs Hamiltonian circuit: cycle traversing every edge ($O(n)$) vs vertex (NP-hard) only once
- MST vs TSP: give paths to (spanning tree, $O(m \log n)$) vs a tour between (NP-hard) all vertices, minimizing the sum of the weights of the edges used
- Class P: Polynomial time problems
- Class NP: Non-deterministic Polynomial
- Class NP-hard: if proving a problem polynomial would mean all NP is polynomial it's NP-hard

7.1 – Reduction

$A < B \rightarrow B$ is used to solve A

$A <_p B \rightarrow A$ reduces to B in polynomial time: a polynomial algorithm exists to convert an input instance for A in one for B that is then used to solve A

if A is NP-hard and $A <_p B \implies B$ is NP-hard

7.2 – NP-hard Problems

- **SAT**: first NP-hard proved, by Cook-Levin theorem
- **3-SAT**: $\text{SAT} <_p \text{3-SAT}$
- **Maximum Independent Set**: $\text{3-SAT} <_p \text{MIS}$ (maximum number of vertices with no edge between them)
- **Hamiltonian circuit**
- **TSP** (Traveling Salesperson Problem): $\text{Hamiltonian circuit} <_p \text{TSP}$
- **Metric TSP**: TSP with triangular inequality on paths (direct paths are always shorter than the ones using other vertices)
- **Maximum clique**: largest complete sub-graph
- **Minimum vertex cover**: minimum number of vertices that “touches” all edges
- **Minimum set cover**: vertex cover $<_p$ set cover (minimum number of subsets to cover an original set)

Lecture 12

8 – Approximation algorithms

8.1 – Vertex cover

Complexity $O(n + m)$

Approximation factor 2

Matching set of edges with no common vertex

8.2 – Metric TSP

Complexity $O(m \log n)$

Approximation factor 2 (tight)

Build an MST with Prim/Kruskal and return the full preorder chain (DFS with pre and post visits (with repetitions)) of the tree

8.2.1 – Eulerian circuit approach

Complexity polynomial

Approximation factor $2/3$

Find a minimum weight perfect matching between odd-degree vertices and add those edges to the MST. Now the graph has all vertices with even degree \Rightarrow it is Eulerian

Return the Eulerian cycle of the graph

A $3/2 - \varepsilon$ approximation has been found, where $\varepsilon = 10^{-36}$

Lecture 17

8.3 – Set cover

Complexity $O(n \cdot |F| \cdot \min\{n, |F|\})$, where $n = |X|$ (cubic)

Approximation factor $\lceil \log_2 n \rceil + 1 = \Theta(\log n)$

Variables:

- X : original set, with all possible elements
- F : set of subsets of X

Greedy algorithm on subset in F with most elements in X . At each step select the subset and remove its elements from X and repeat

9 – Randomized algorithms

- Las Vegas: always correct (randomized quicksort)
- Monte Carlo: may return wrong values, though high probability of correct result
 - One sided: decision problems give only false positives/negatives
 - Two sided: decision problems may fail in any case

Lecture 19

High probability algorithm A_Π for problem Π

has complexity $f(n)$ / is correct

with high probability if

$$\exists c, d > 0. \Pr (A \text{ has complexity} > cf(n)) / \Pr (A \text{ is not correct}) < \frac{1}{n^d}$$

10 – Minimum cut – Karger

Complexity $O(n^4 \log n)$

Minimum number of edges to remove, in order to disconnect the (multi)graph

10.1 – Algorithm

Repeat *Full Contraction* k times, to reduce error

Karger returns the minimum with *high probability* ($\Pr (\text{fail}) < \frac{1}{n^d}$) with

$$k = \frac{dn^2 \ln n}{2} = \Theta(n^2 \log n)$$

10.2 – Definitions

10.2.1 – Multigraphs

Multiplicity $m : \mathbb{S} \rightarrow \mathbb{N}, m(e) = \text{occurrences of an element } e \in \text{multiset } \mathbb{S}$

$\mathcal{G} = (\mathcal{V}, \mathcal{E})$ is a multigraph, where \mathcal{E} is a multiset

10.2.2 – Full Contraction

Complexity $O(n^2)$

Choose a random edge and contract on it, until two vertices remain

Contraction contract a graph \mathcal{G} on edge $(u, v) \in \mathcal{E}$ (join vertices of the edge):

- Delete u
- Delete all edges between u and v
- Move all edges of u to v

10.3 – Karger-Stein

Complexity $O(n^2 \log^3 n)$

Avoids first $\frac{n}{\sqrt{2}}$ iterations

10.4 – 2020 version

Complexity $O(m \log n)$