

# Real-Time Kernels & Systems

---

**Academic Year 2022/2023**  
**Master Degree in Computer Science**  
**Department of Mathematics**  
**University of Padua**  
**Tullio Vardanega, [tullio.vardanega@unipd.it](mailto:tullio.vardanega@unipd.it)**

# Lecture topics

## Bibliography

- J. Liu, "Real-Time Systems", Prentice Hall, 2000
- A. Burns and A. Wellings, "Analysable Real-Time Systems" Amazon Books, 2016
- State-of-the-art literature

1. Introduction
2. Scheduling basics
3. Fixed-priority scheduling
  - a. Basic workload models
  - b. Task interactions and blocking
  - c. Exercises and further model extensions
4. Implementation issues
  - a. Programming real-time systems (in Ada)
  - b. A look under the hood
5. Distributed systems
6. Execution-time analysis
7. Multicore systems
  - a. Initial reckoning
  - b. Seeking the lost optimality
  - c. Global resource sharing
8. Mixed-criticality systems
9. Predictable parallel programming

# Protocol

## ■ Learning outcomes

- Realize the existence and the needs of software systems whose response time is critical to their use and consequently to their design
- Understand the principles, methods, techniques and technology required to develop them so as to guarantee predictability

## ■ Instructional methods

- Web resources in complement to slide desks and presentations
- Reciprocal feedback
  - Instructor to student, about incremental (self-)assignments
  - Student to instructor, about the progression of learning
- Interaction
  - During class
  - Via posts in the Moodle's class billboard, for all that must be public
  - Email otherwise

# 1. Introduction

---

**Where we make some initial acquaintance  
with what real-time systems are and why  
they came about, and then take a first look  
at their abstract concept**

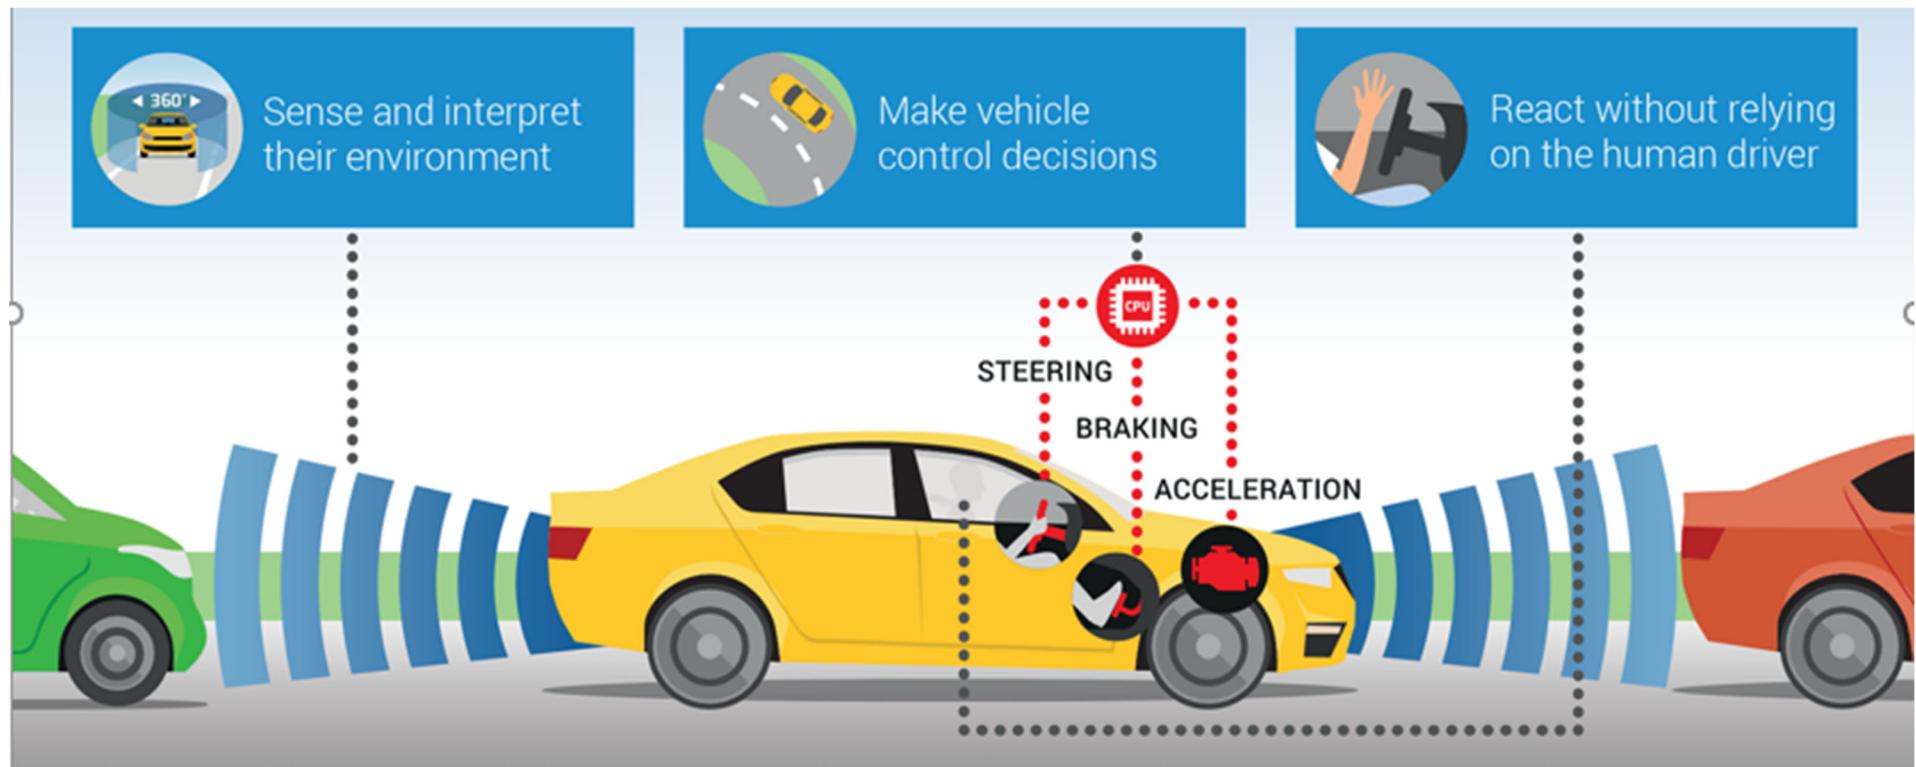
# Initial intuition – 1/2

## ■ Real-time system

- An aggregate of computers, I/O devices and *application-specific software*, characterized by
  - Continuous interaction with the external environment
  - To control it after mission-specific goals
  - Capturing variations of the state of the system in relation to its physical environment, and reacting to them in a timely fashion
- Comprised of system activities subject to timing constraints
  - Reactivity, accuracy, duration, completion, responsiveness
  - All dimensions of *timeliness*
- System activities inherently *concurrent* and increasingly *parallel*
- The satisfaction of all system constraints must be proved

# Illustrative example

- ... capturing variations of the state of the system in relation to its physical environment, and reacting to them in a timely fashion



# Initial intuition – 2/2

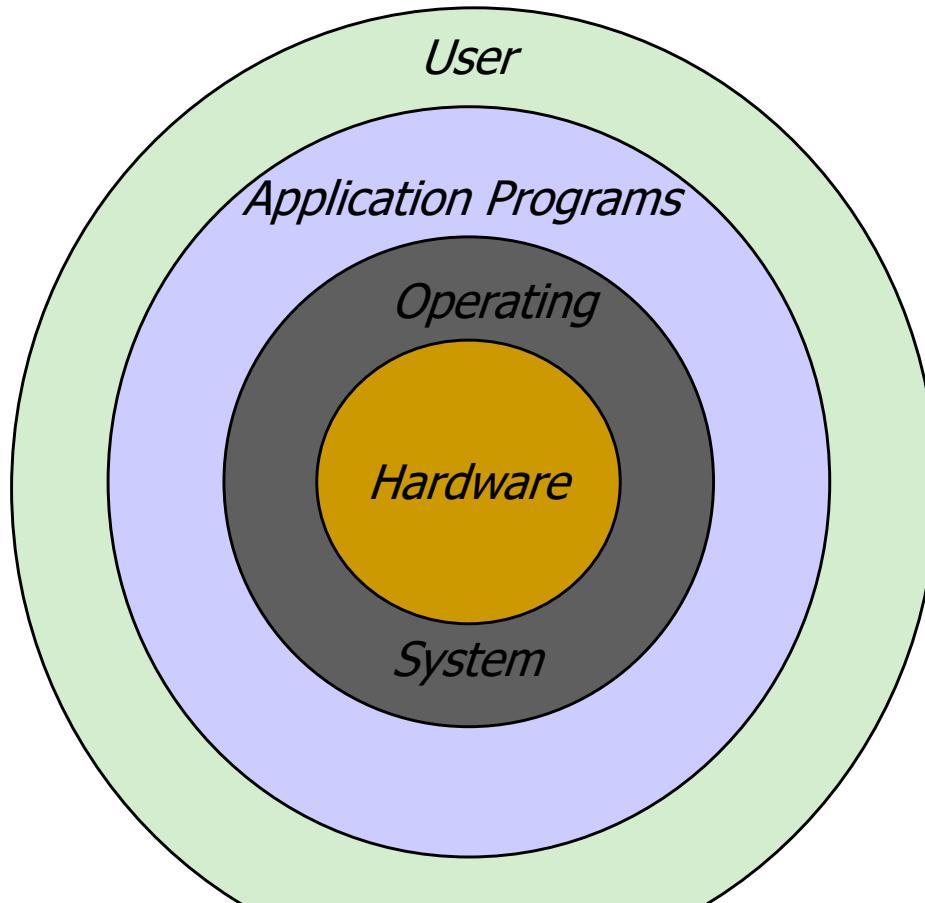
## ■ Real-time system (cont'd)

- Operational correctness does not depend solely on the logical result but also on the time at which the result is produced
  - The computed response has an application-specific utility
  - Correctness is defined in the value domain *and* in the time domain
  - A logically-correct response produced later than due may be bad

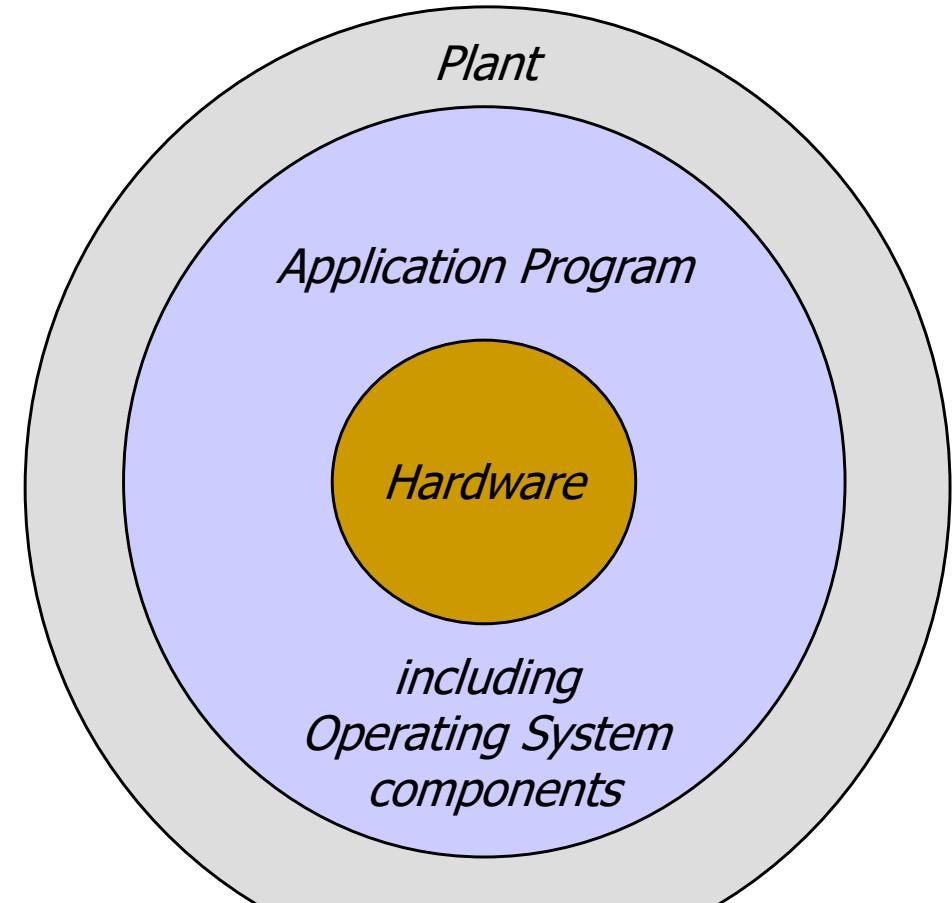
## ■ Embedded system

- The computer and its software are fully immersed in an *engineering system* comprised of the external environment subject to its control
  - Defined in terms of physical attributes that can be interacted with

# Embedded system – 1/2



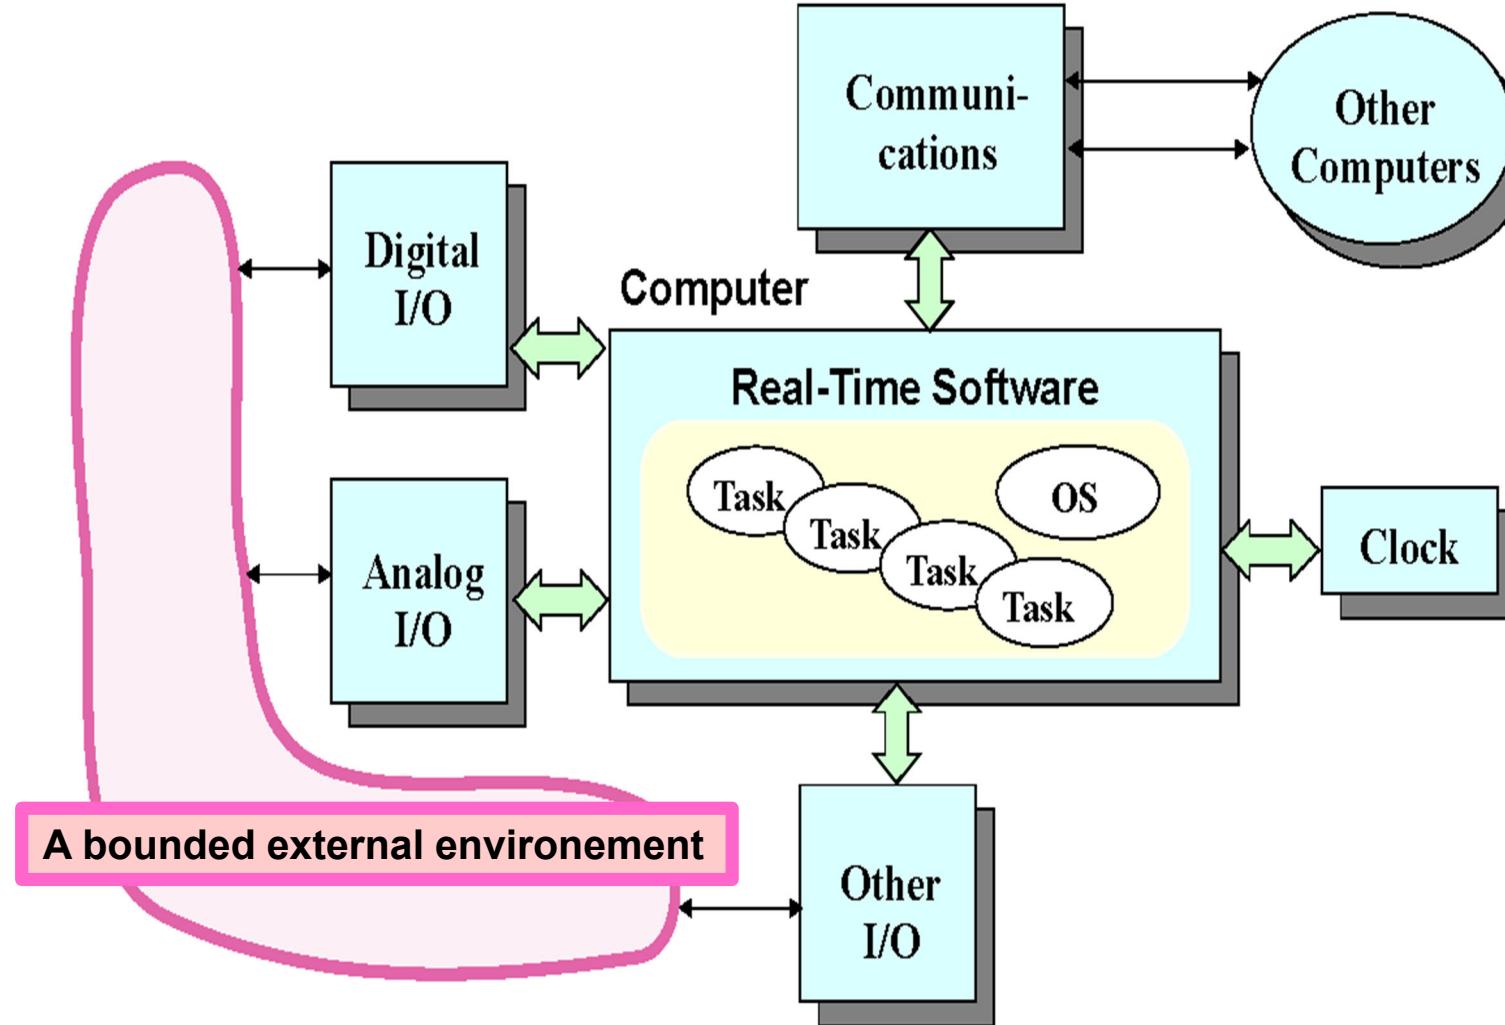
**Typical General-Purpose Computing Configuration**



**Typical Embedded Computing Configuration**



# Embedded system – 2/2



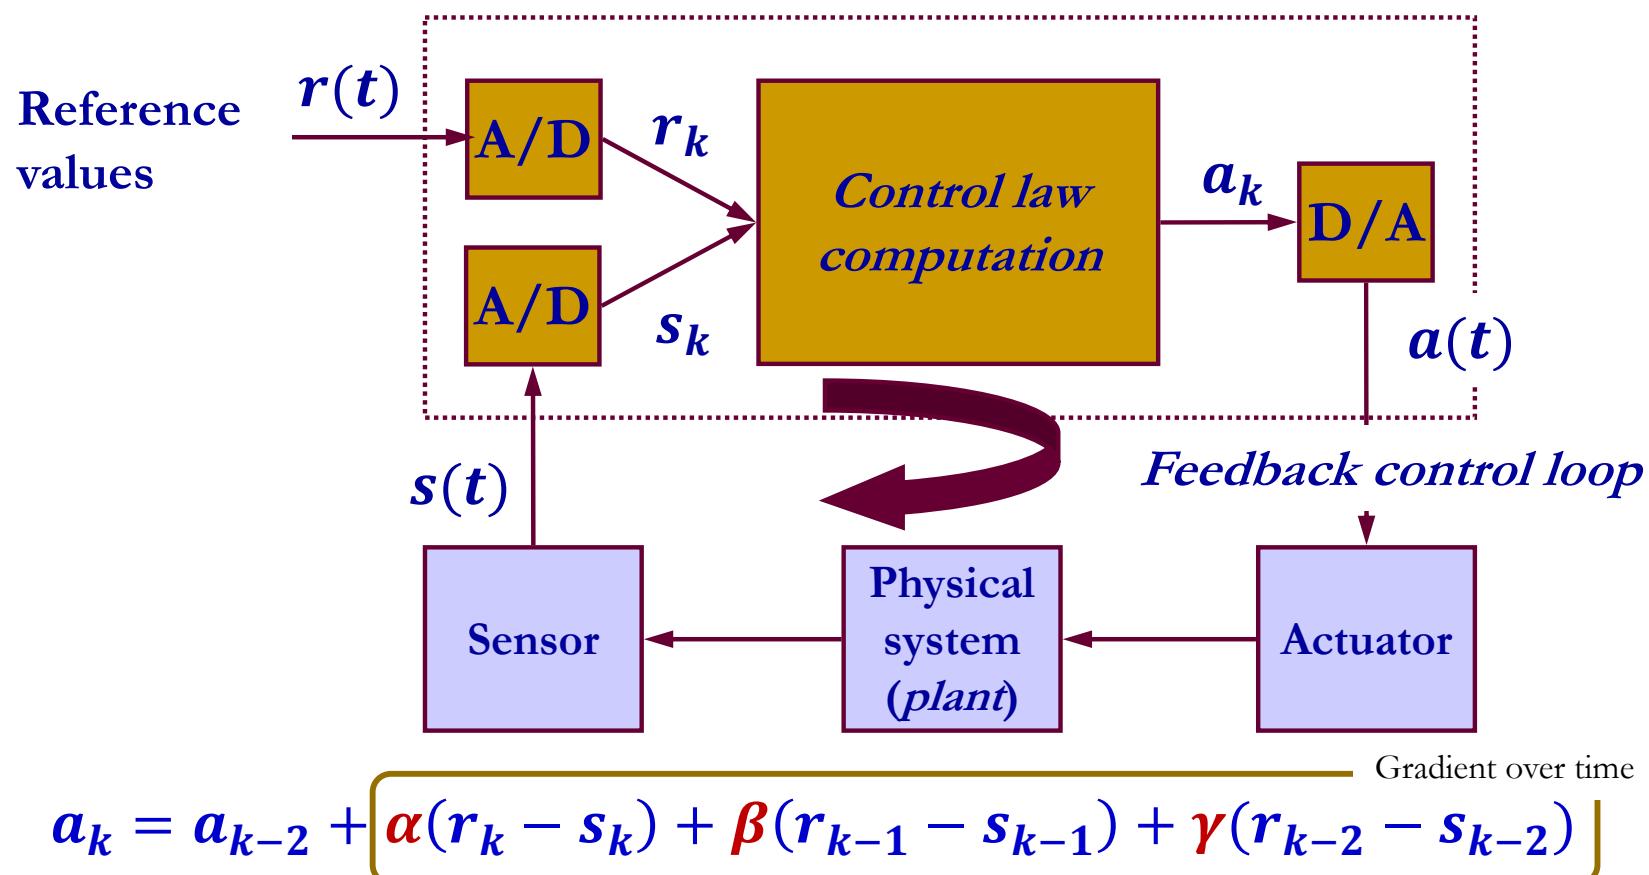
# Cybernetics



- Born in 1948 as *the science of control systems*, prerequisite to the **automation of control**
- From Greek's *κυβερνητής* (Latin's “governator”), helmsman
  - *Sensing* the external (physical) environment
  - *Computing* the distance from the expected status
  - *Actuating* devices to effect the system or the environment, to stay in course
- Every actuation performed on the external environment causes *feedback* that must be accounted for
- The goal of cybernetics is to calibrate actuation so that the system goal is reached with bounded feedback

# Automation of control – 1/4

- A digital system comprised of sensors and actuators

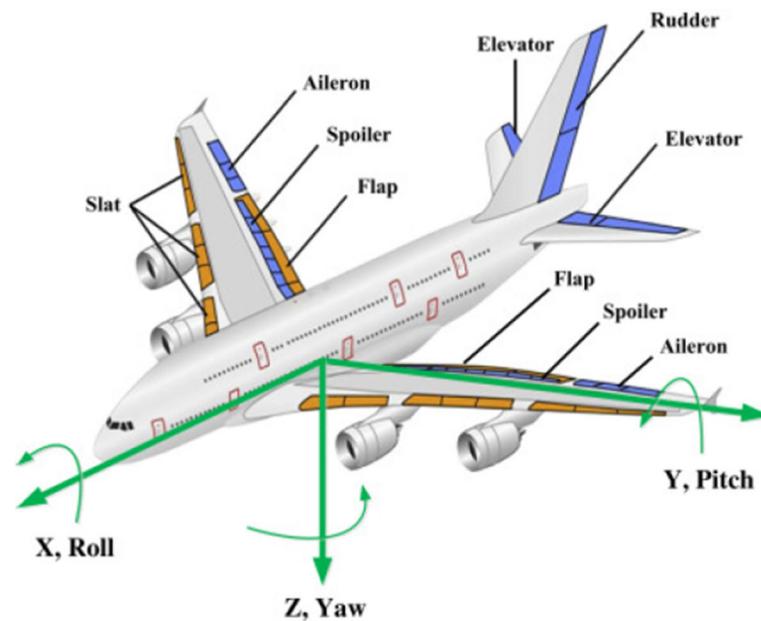


# Automation of control – 2/4

## ■ Factors of influence

- Quality of response (*responsiveness*)
  - Sensor sampling is typically periodic with period  $T$ 
    - For the convenience of control theory
  - Actuator commanding is produced at the time of the next sampling
    - As part of feedback control mathematics
  - System stability degrades with the width of the sampling period
- Plant *capacity*
  - Good-quality control reduces oscillations
  - A system that needs to react rapidly to environmental changes and is capable of it within *rise time*  $R$  requires higher frequency of actuation and thus faster sampling → hence shorter  $T$
  - A rule-of-thumb  $R/T$  ratio normally ranges [10 .. 20]

# Automation of control – 3/4



Any three-dimensional rotation can be described as a sequence of *roll* ( $x$ ), *pitch* ( $y$ ), *yaw* ( $z$ ) rotations (Euler angles)

- Complex systems must support multiple distinct periods  $T_i$ 
  - A **harmonic** relation among all  $T_i$  does help
    - This removes the need for concurrency of execution in the relevant computations
    - But it causes coupling between possibly unrelated control actions which is a poor architectural choice
  - There may be diverse components of speed
    - *Forward, side slip, altitude*
  - As well as diverse components of rotation
    - *Roll, pitch, yaw*
  - Each of them requires separate control activities each performed at a specific rate

# Automation of control – 4/4

## (Artificially) harmonic multi-rate system

- 180 Hz cycle
  - Check all sensor data and select sources to sample
  - Reconfigure system in case of read error
- 90 Hz cycle
  - Perform control law for pitch, roll, yaw (internal loop)
  - Command actuators
  - Perform sanity check
- 30 Hz cycle
  - Perform control law for pitch, roll, yaw (external loop) and integration
- 30 Hz cycle
  - Capture operator keyboard input and choice of operation model
  - Normalize sensor data and transform coordinates; update reference data
- *Can you figure how those activities can progress together?*

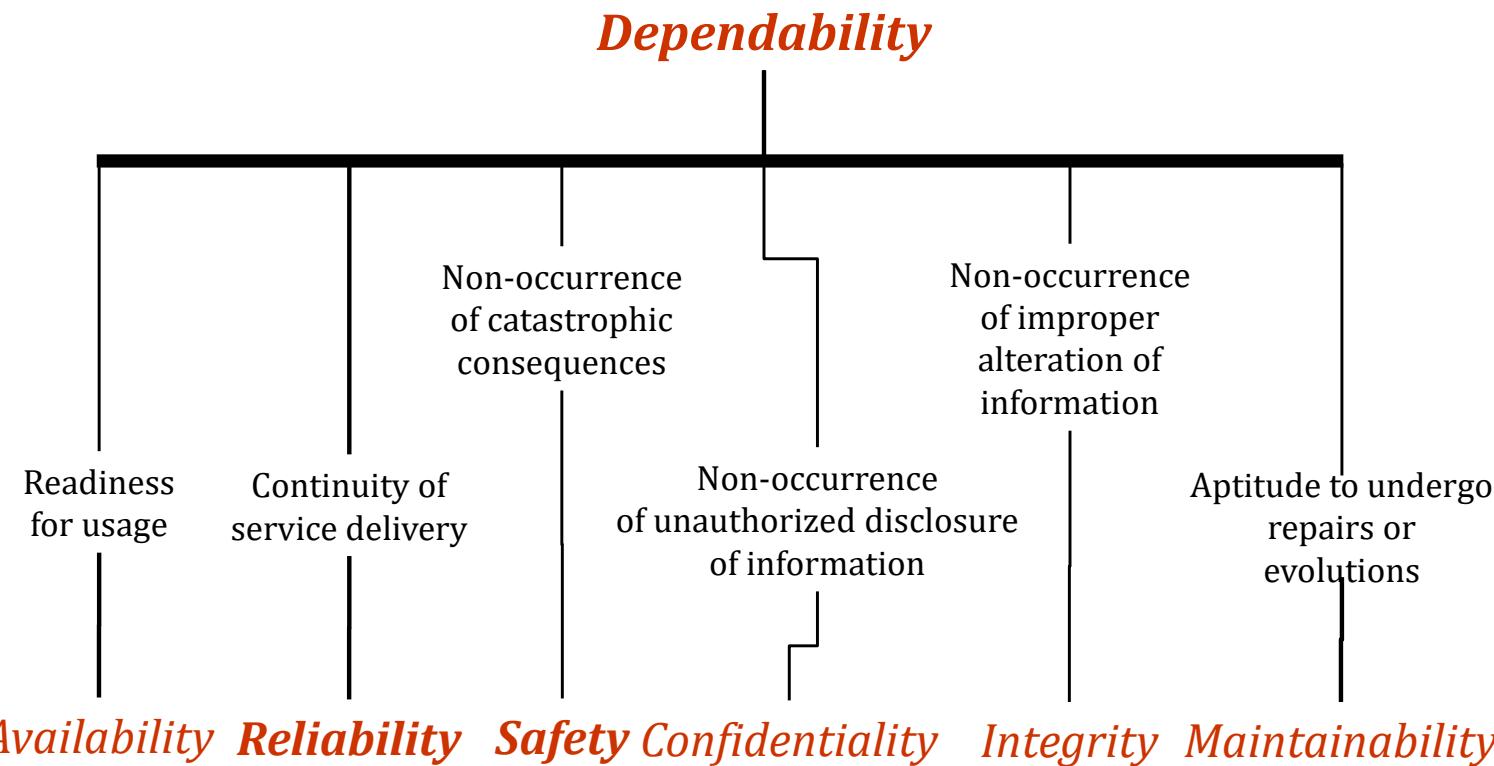
# Application requirements

- A control system consists of (distributed) logical and physical resources governed by a *real-time operating system*, RTOS
- The system design overall must meet stringent ***reliability*** requirements
  - Measured in terms of *maximum acceptable probability of failure*
  - For example:  $10^{-9}$ /hour of flight for the Airbus A-3X0 control system
    - One failure allowed in  $10^9$  hours of flight ( $> 114k$  years!)

# Traits of real-time systems – 1/2

- **Complex**
  - In algorithms, because of the need to apply discrete control over analog and continuous physical phenomena
  - In development, owing to highly taxing verification and validation processes
- **Heterogeneous** for components and processing activities
  - Multi-disciplinary engineering (spanning control, SW, and system)
- Extremely **variable** in size and scope
  - From tiny and pervasive (nanodevices) to very large (aircraft, plant)
  - In all cases, finite in computational resources
- Provably ***dependable*** [glossary]

# Dependability attributes



- See the **self-study material** posted on the lecture schedule page

# Traits of real-time systems – 2/2

- Respond to events triggered by the external environment as well as by the passing of time
  - Their schedule is both *event-driven* and *time-driven*
- Ensure continuity of operation
  - Operating without (constant) human supervision
- Inherently *concurrent* and increasingly *parallel* [glossary]
- Temporally ***predictable*** [glossary]
  - Need for static (off-line) verification of correct temporal behaviour
  - How does that relate to ***determinism***? [glossary]

# Concurrency vs. parallelism

- **Concurrent** programming allows using multiple logical threads of control to reflect *cohesively* the collaborative structure of the solution
  - I do “my” specialized bit, you do “yours”; our mutual waiting is not wasteful
  - Threads form the architecture: they are *long-lived*
- **Parallel** programming promotes a divide-and-conquer logic to solve a problem, with multiple threads that work *independently* on the problem space
  - I work as fast as I can and know nothing about you
  - Threads are mindful of throughput: they are *short-lived*

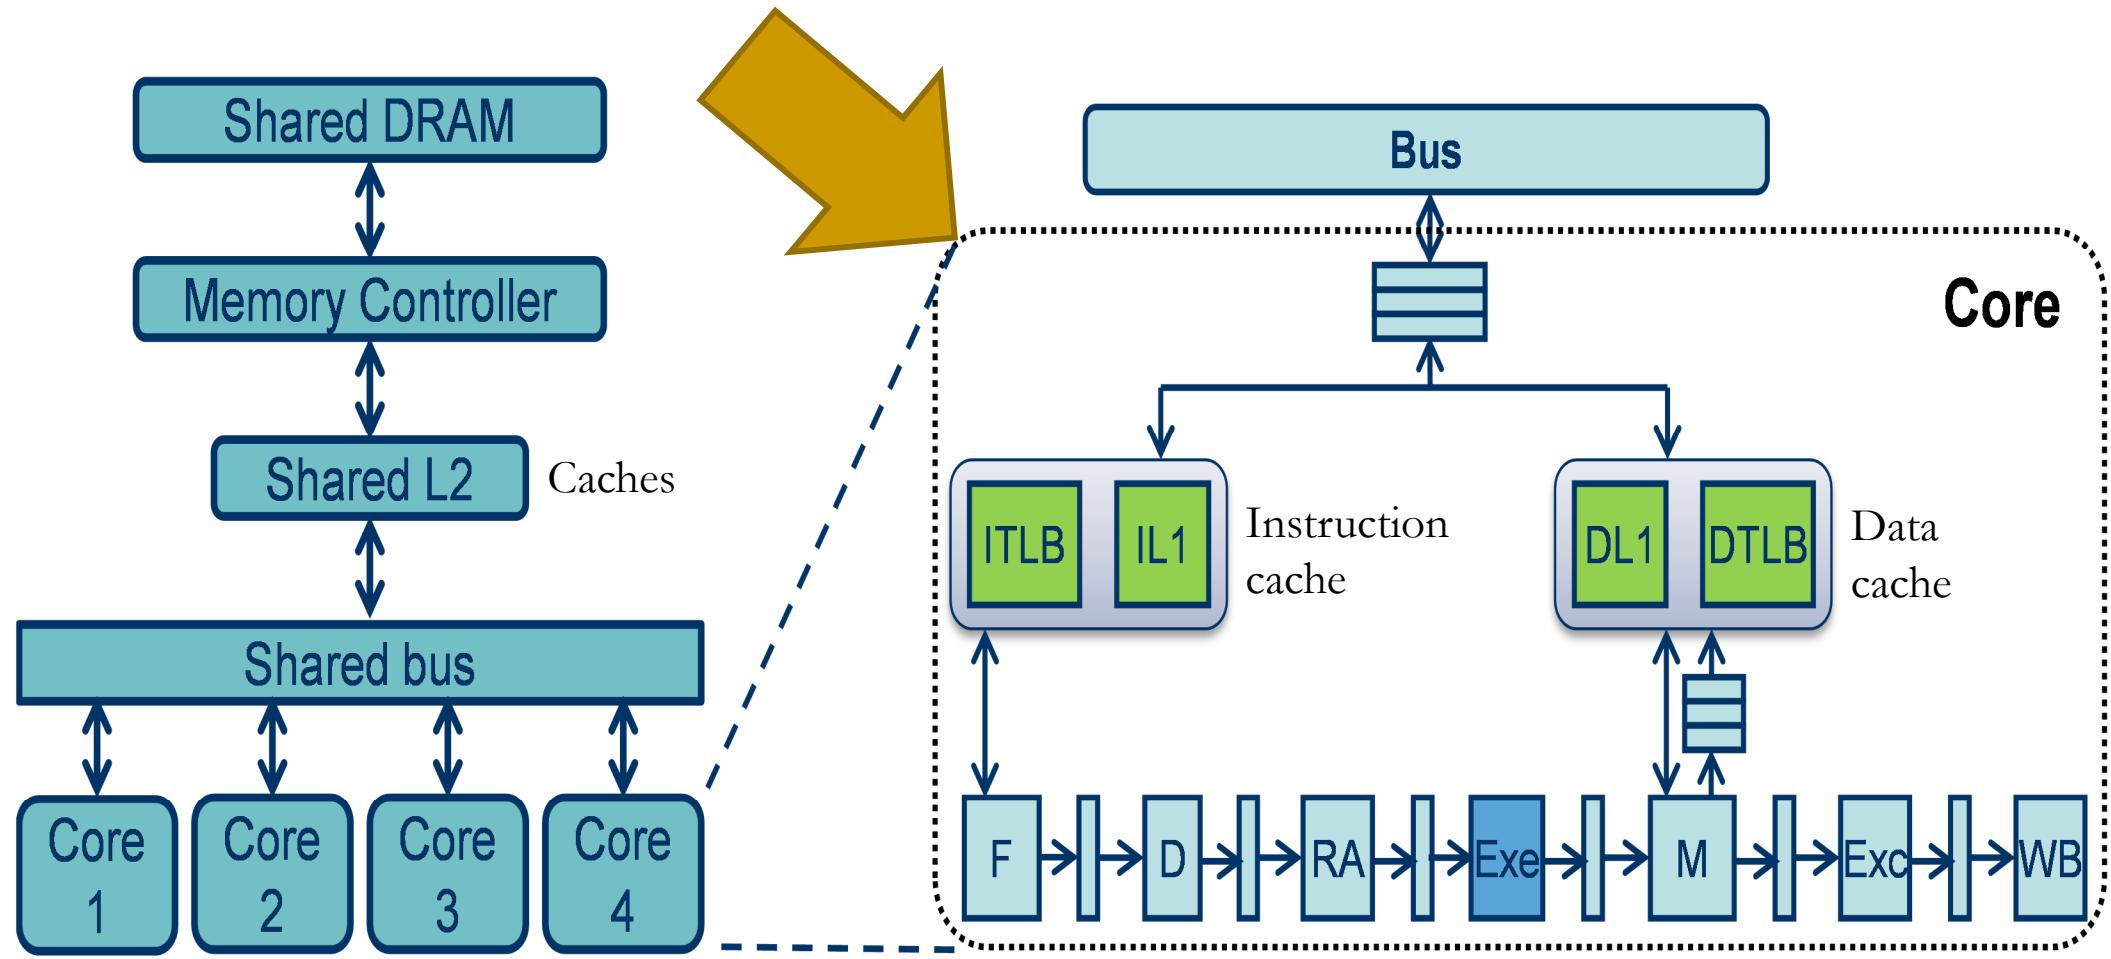
# Predictability vs. determinism

- *Predictability* is what can be soundly established a priori
- It may be regarded as a continuum
- Its highest end is full a-priori knowledge
  - Which allows for deterministic reasoning, and yields absolute certainty
- Its lowest end is total absence of a-priori knowledge
  - In the style of “See what happens ...”
- Seeking predictability implies reasoning about kinds and degrees of uncertainty, pursuing an acceptable balance
- We must reason conservatively
- Considering *worst-case* conditions, which may be difficult to capture, as very rarely we have full a-priori knowledge of the factors of influence
- At the same time, we must avoid exceedingly pessimistic reasoning

# Meeting real-time requirements

- Minimizing the application tasks' average response time matters to general-purpose computing, but it does **not** to RTS!
- *Real-time* computing is **not** equivalent to *fast* computing
- Given real-time requirements and a concrete implementation, how can one show that those real-time requirements are met?
  - Testing and simulation are **not** sufficient
  - Telling example: maiden flight of space shuttle, April 12, 1981:  
there was a theoretical  $\frac{1}{67}$  probability of a *transient overload* occurring at system initialization
  - It never showed up in testing; it did at launch
- The answer to that lies in seeking system-level ***predictability***
  - Which requires understanding the *worst case*
  - Which further requires understanding how execution works ...

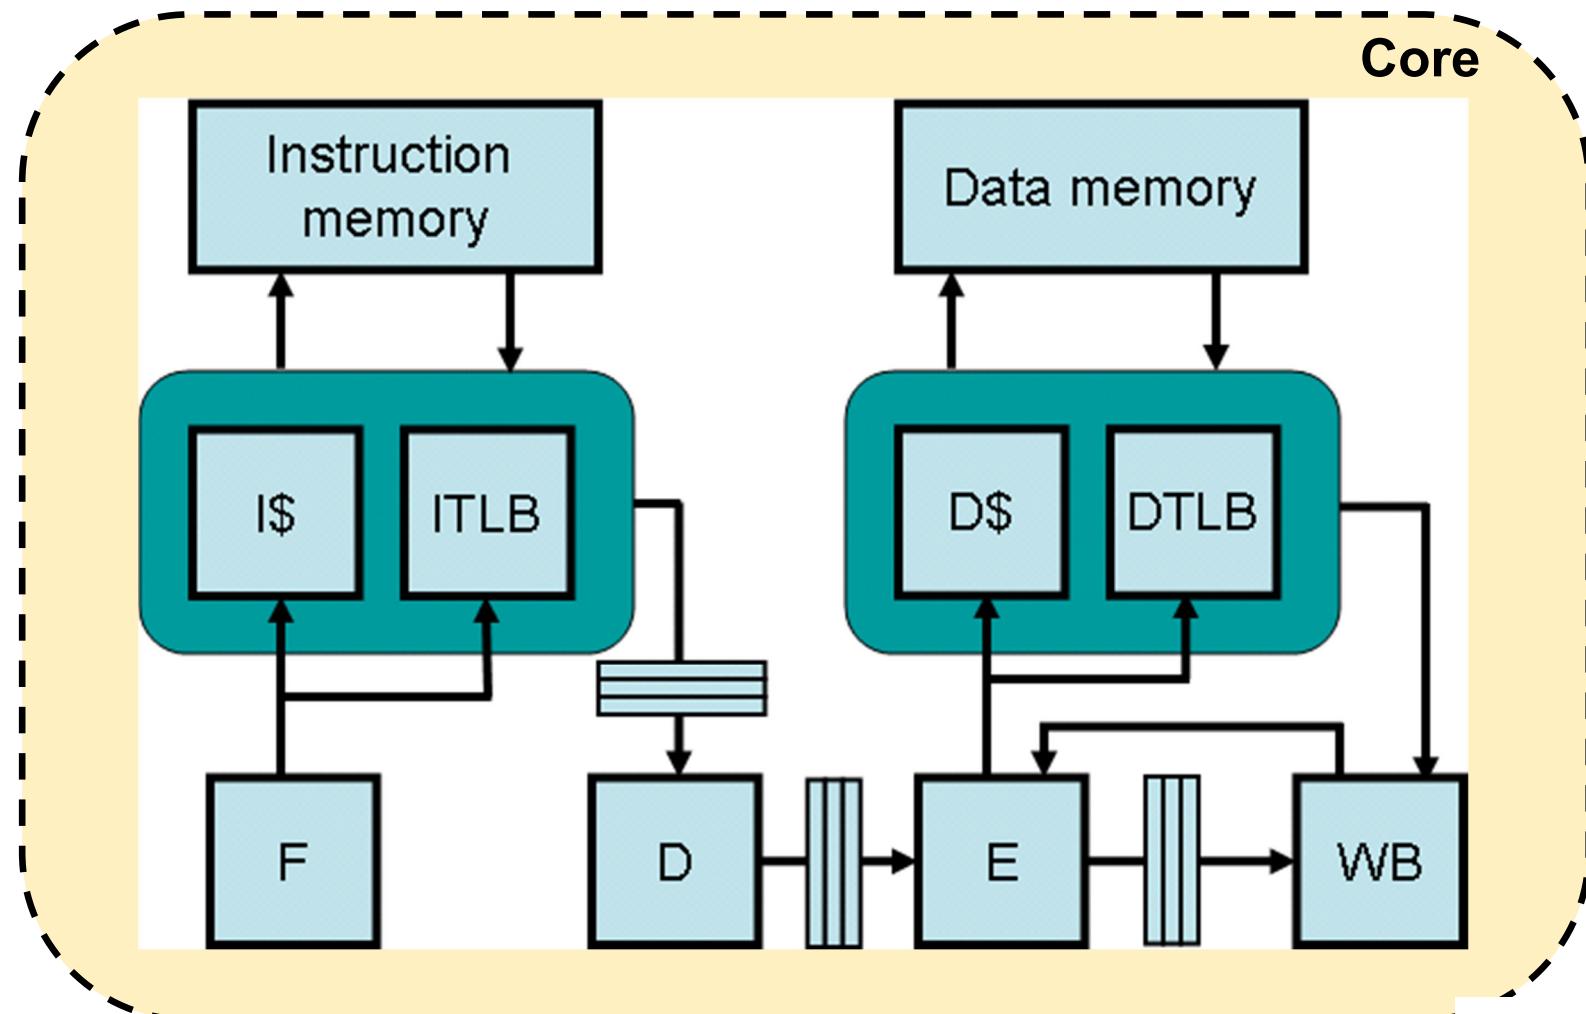
# Understanding the processor – 1/2



Courtesy of

**PROXIMA**

# Understanding the processor – 2/2



Courtesy of

**PROXIMA**

# Definitions – 1/6

## ■ *Task*

- Concurrent unit of functional architecture: they *never* end
- Issues *one job at a time*, until completion, to perform actual work
- Real-time systems' tasks are typically *recurrent*:
- The body of a task may be seen as an *endless loop*

## ■ *Job*

- Unit of work called into execution by a task, following a given law of activation
- Competes for CPU time under control by a system-level scheduler
- Needs physical and logical *resources* to execute
- The job may be seen as a top-level procedure, which carries out the actual functional work required of the corresponding task and then terminates

# Definitions – 2/6

## ■ *Release time*

- When it occurs, one job becomes eligible for execution
- The corresponding trigger is called *release event*
  - There may be some temporal delay between the arrival of the release event and when the scheduler recognizes the job as ready
- It may be set at some *offset* from system start time
  - For example, to avoid congestion on access to the CPU
  - The offset of the *first* job of task  $\tau$  to the system start time is named *phase*,  $\varphi$ , and it is one of the attributes of  $\tau$

# Definitions – 3/6

## ■ *Jitter*

- Variability in the arrival of the release time
- Variability in the arrival of input (data freshness)
- Variability in the delivery of output (stability of control)

## ■ *Inter-arrival time*

- Separation between release time of any two successive jobs
- Job is *periodic* if inter-arrival time is constant
- Job is *sporadic* if a guaranteed minimum such value exists
- Job is *aperiodic* otherwise

## ■ *Execution time, C*

- For any job  $J_i$ ,  $C_i$  may vary between a *best-case* (**BCET**)  $C_i^b$  and a *worst-case* (**WCET**)  $C_i^w$

# Definitions – 4/6

## ■ ***Deadline***

- The latest time by which a job must complete its execution
- May be < (*constrained*), = (*implicit*), > (*arbitrary*) than the next job's release time

## ■ ***Response time, R***

- The time span between the job's release and its actual completion
- It may be longer than the job's execution time ( $R \geq C$ )

## ■ The longest admissible response time for a job $j_i$ is termed the job's *relative deadline*, $D_i$

- $R_i \leq D_i$  should always hold

## ■ The algebraic summation of release time and relative deadline is termed *absolute deadline*, $d_i$ , which is a specific point in time

# Definitions – 5/6

- Deadlines are said to be **hard**
  - If the consequences of a job completing past its deadline are serious and possibly intolerable
  - Satisfaction of all deadlines must be proven offline
- Deadlines are said to be **soft**
  - If the consequences of a job *occasionally* completing past the assigned deadline are tolerable
  - The quantitative interpretation of “occasional” may be established in probabilistic terms or in terms of **utility**
- Deadlines are said to be **firm**
  - When they are soft but have utility  $\leq 0$  past the deadline point, and therefore may cause damage if missed

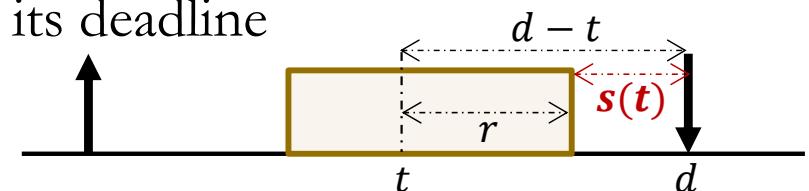
# Definitions – 6/6

## ■ **Laxity**, aka **slack** at time $t$

- $s(t) = (d - t) - r$  is the *slack* at time  $t$  of job  $J$  with absolute deadline  $d$  and remaining time of execution  $r$ 
  - A job with non-negative laxity meets its deadline

## ■ **Tardiness**

- The span between a job's response time and its deadline
  - A job with negative laxity has tardiness

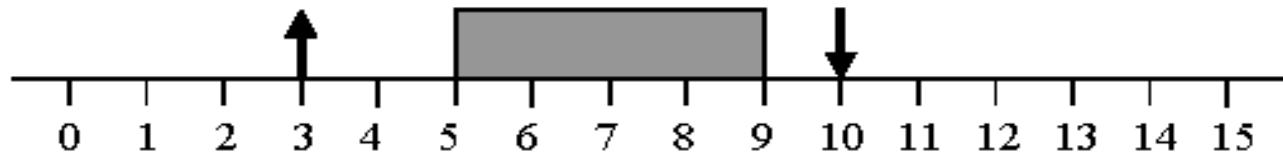


## ■ **Usefulness**

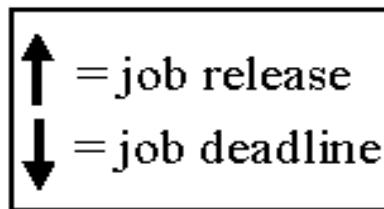
- Value of (residual) utility of the job's computational product as a function of its tardiness

# An example timeline

## Example



Without loss of generality,  
timelines have events occur  
at integral points in time

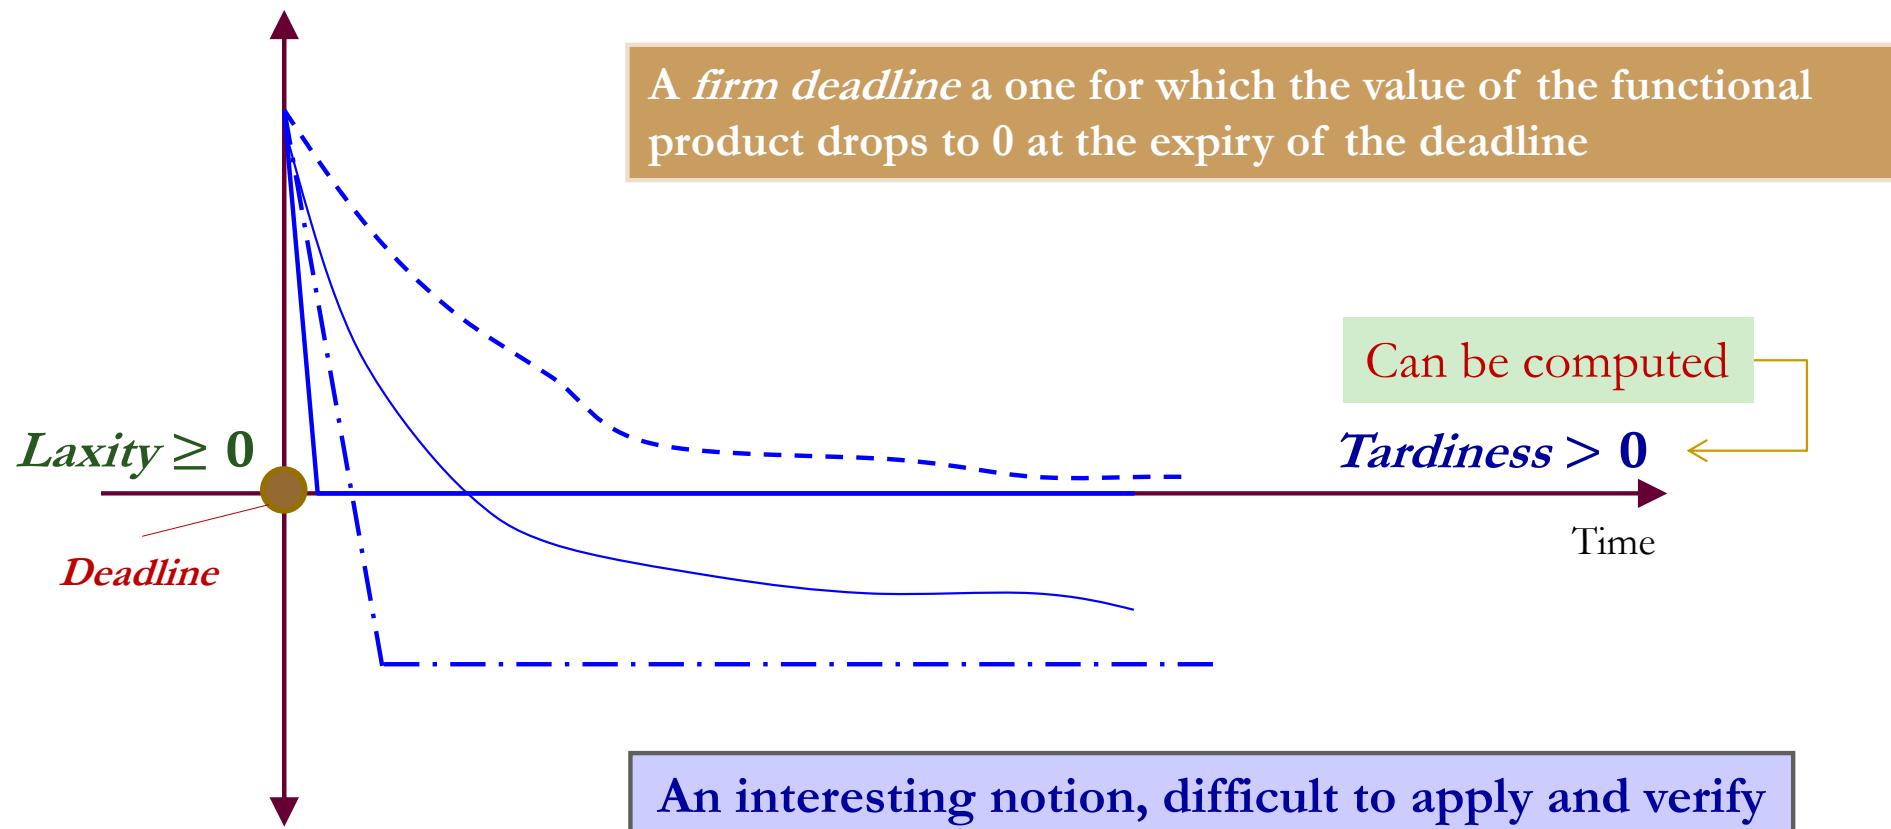


Job is released at time 3.  
It's (absolute) deadline is at time 10.  
It's relative deadline is 7.  
It's response time is 6.

# Utility function

Difficult to quantify

→ **Usefulness**



# An initial taxonomy

## ■ *Periodic* tasks

- Their jobs become ready at regular intervals of time,  $T$
- Their arrival is synchronous to some time reference

## ■ *Aperiodic* tasks

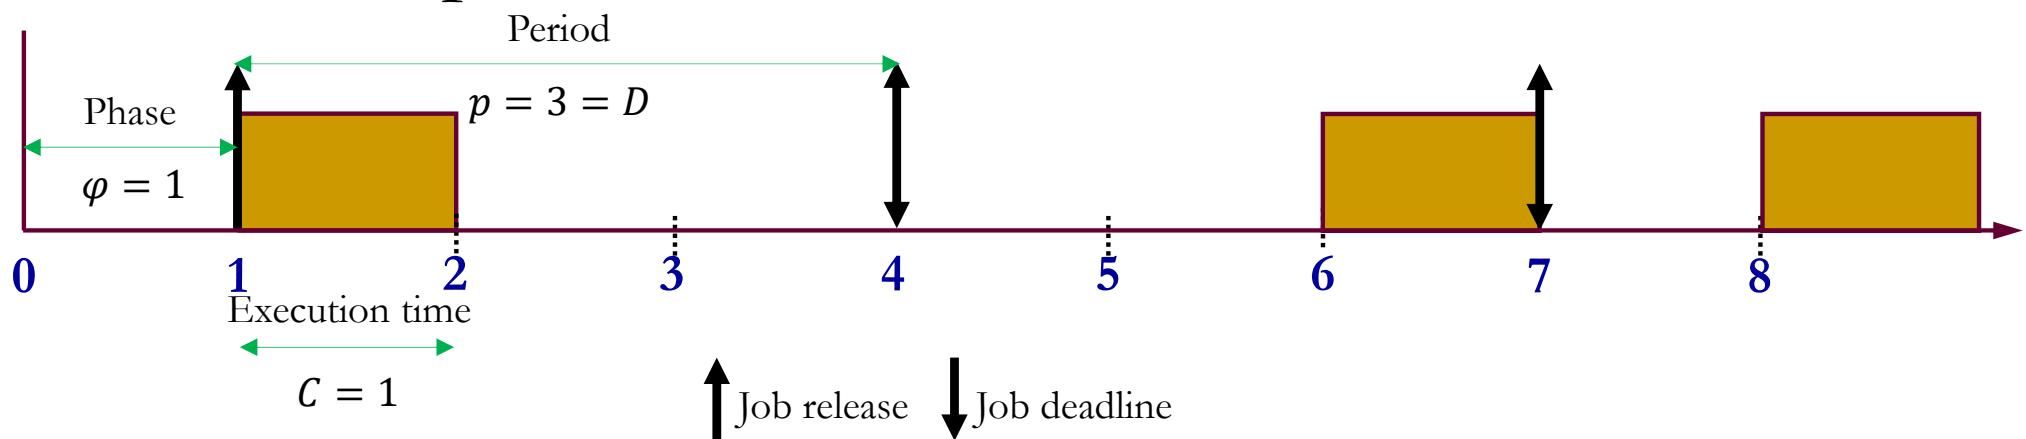
- Recurrent but irregular
- Their arrival cannot be anticipated (asynchronous)

## ■ *Sporadic* tasks

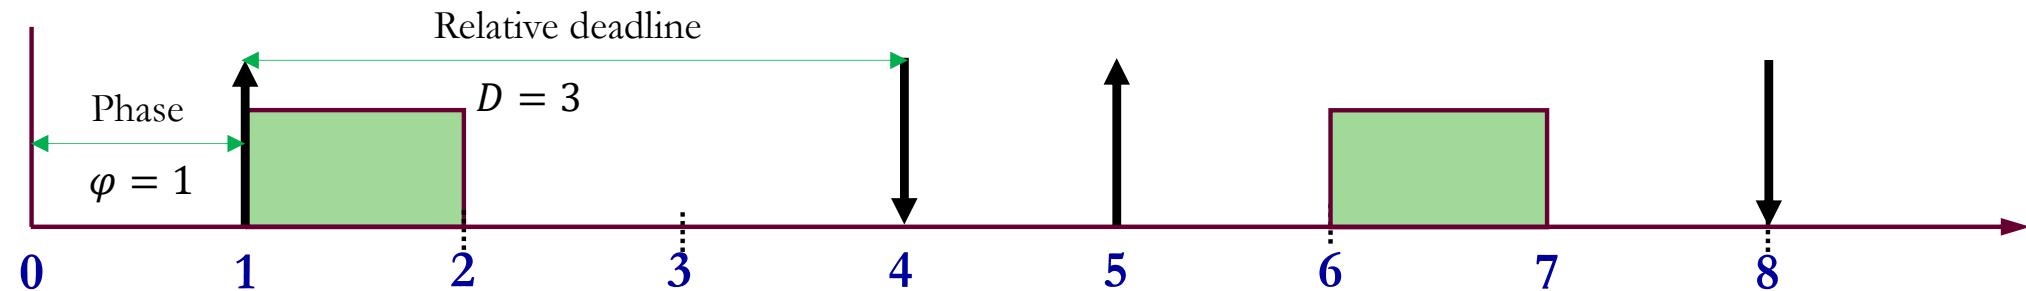
- Their jobs become ready at variable times but at bounded minimum distance from one another

# Example

- This is a *periodic* task



- This is a *sporadic* task



# Abstract modelling – 1/8

## ■ *Active resources* (processor, server)

- They “do” what they have to do
  - Execute machine instructions, move data, process queries, etc.
- Jobs must acquire them to progress toward completion
  - Contention occurs on access to them

## ■ Active resources have a type

- Those of the same type can be used interchangeably
- Those of different types cannot
  - For example, processors may have different speed, which affects the rate of progress for the jobs that run on them

# Abstract modelling – 2/8

## ■ ***Passive resources*** (memory, shared data, semaphores, ...)

- A passive resource doesn't do anything per se, but jobs *may need* it to make progress
- They may be reused if use does not exhaust them
  - If always available in sufficient quantity to satisfy all needs, they are said to be *plentiful* and can be ignored
- Passive resources that matter to real-time systems are those that may cause *bottlenecks*
  - Access to memory may matter more (owing to *arbitration*) than memory itself (which may be considered plentiful)

# Abstract modelling – 3/8

- Permissibility of job **preemption**
  - May depend on the capabilities of the execution environment (e.g., *non-reentrancy*) but also on the programming style
  - Preemption causes time and space overhead
- Job **criticality**
  - Akin to a criterion of execution eligibility
  - It indicates which activities must be guaranteed no matter what
- Permissibility of resource preemption
  - Some resources are intrinsically preemptable
  - Others do not permit it
  - Which ones?



# Abstract modelling – 4/8

## ■ Fixing execution parameters

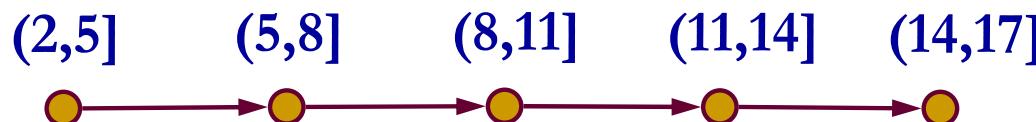
- The time that elapses between when a periodic job becomes ready and the next period  $p$  is certainly  $\leq p$
- Setting phase  $\varphi > 0$  and relative deadline  $D < p$  for a job may help limit its output jitter (**why?**)
- The jobs of a system may be independent of one another
  - They may execute in any relative order
- Or they may be subject to *precedence constraints*
  - As is the case in collaborative architectural styles
  - E.g., producer – consumer

# Task precedence graphs

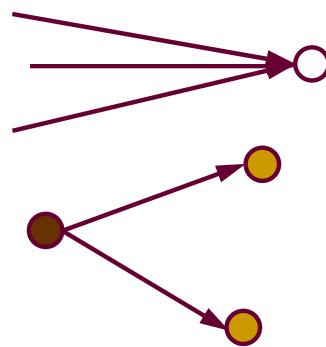
Relative deadline  
Phase      Period = 2



Independent jobs



Dependent jobs



Job of type AND (join)

Job of type OR (branch)  
typically followed by  
a join job

# Types of precedence constraints

- The release time of a job with successors cannot follow the release of any of its successors
- ***Effective release time (ERT)***
  - For a job  $J_i$  with predecessors  $\{J_{k=1,\dots,i-1}\}$ ,  $ERT_i$  is the *latest* value between its own release time and the maximum ERT of its predecessors,  $ERT_k$ , plus  $C_k$
  - $J_i$  cannot be released before  $\max_k(R_k)$ , when all its predecessors have completed
- The absolute deadline of a job with predecessors cannot precede the absolute deadline of any of its predecessors
- ***Effective deadline (ED)***
  - For a job  $J_i$  with successors  $\{J_{k=i+1,\dots,n}\}$ ,  $ED_i$  is the *earliest* value between  $D_i$  and the minimum ED of its successors,  $ED_k$ , less  $C_k$
  - $D_i \leq \min_k(ED_k - C_k)$ : no successor may run before the latest end of its predecessors
- For single processors with preemptive scheduling, ERT and ED are the only precedence constraints that matter

# Abstract modelling – 5/8

## ■ *Periodic model*

- Comprises periodic *and* sporadic jobs
- Accuracy of representation decreases with increasing jitter and variability of execution time
- **Hyperperiod**  $H_S$  of task set  $S = \{\tau_i\}, i = 1, \dots, N$ 
  - Defined as LCM (least common multiple) of task periods  $\{p_i\}$
  - It is the time span within which all tasks have issued at least one job
- **Utilization**
  - For every task  $\tau_i$  : defined as the ratio between execution time and period :  $U_i = \frac{c_i}{p_i} \leq 1$
  - For the system (*total utilization*) :  $U = \sum_i U_i \leq m$ , where  $m$  is the number of CPUs ( $m = 1$ , for now)

# Abstract modelling – 6/8

- Selecting jobs for execution
  - The scheduler assigns a job to the processor's CPU
  - The resulting assignment is termed ***schedule***
- A schedule is ***valid*** if
  - Each processor is assigned to at most 1 job at a time
  - Each job is assigned to at most 1 processor at a time
  - No job is scheduled before its release time
  - The scheduling algorithm ensures that the amount of processor time assigned to a job is  $\geq$  than its BCET and  $\leq$  than its WCET
  - All precedence constraints in place among tasks as well as among resources are satisfied

**BCET**: best-case execution time

**WCET**: worst-case execution

# Abstract modelling – 7/8

- A *valid schedule* is said to be **feasible** if it satisfies the temporal constraints of every job
- A *job set* is said to be **schedulable** by a scheduling algorithm if that algorithm always produces a *valid* schedule for that problem
- A *scheduling algorithm* is **optimal** if it always produces a *feasible* schedule when one exists
- Actual systems may include multiple schedulers that operate in some hierarchical fashion
  - E.g., some scheduler governs access to logical resources; some other schedulers govern access to physical resources

# Abstract modelling – 8/8

- Two algorithms are of prime interests for real-time systems
  - The *scheduling algorithm*, which we should like to be *optimal*
    - Comparatively easy problem
  - The *analysis algorithm* that tests the *feasibility* of applying a scheduling algorithm to a given job set
    - Much harder problem
- The scientific community, but not always consistently, divides the analysis algorithms in
  - **Feasibility tests**, which are exact (necessary and sufficient)
  - **Schedulability tests**, which are only sufficient

# Cute teasers

- Björn Brandenburg's post on SIGBED Blog ...
  - <https://sigbed.org/2020/09/05/liu-and-layland-and-linux-a-blueprint-for-proper-real-time-tasks/>
- ... shows a cute way to reconcile the basic theory seen so far and simple-but-sound conveniency programming
  - Check it out
- Linux is certainly not what you would expect to find in a real-world embedded real-time system
  - Yet, lots of users like “convenience-development”
  - <https://www.zdnet.com/article/to-infinity-and-beyond-linux-and-open-source-goes-to-mars/>



# Further characterization – 1/2

	Time-Share Systems	Real-Time Systems
Capacity	High throughput	Ability to meet timing requirements: Schedulability
Responsiveness	Fast average response	Ensured worst-case latency
Overload	Fairness	Stability of critical part



# Further characterization – 2/2

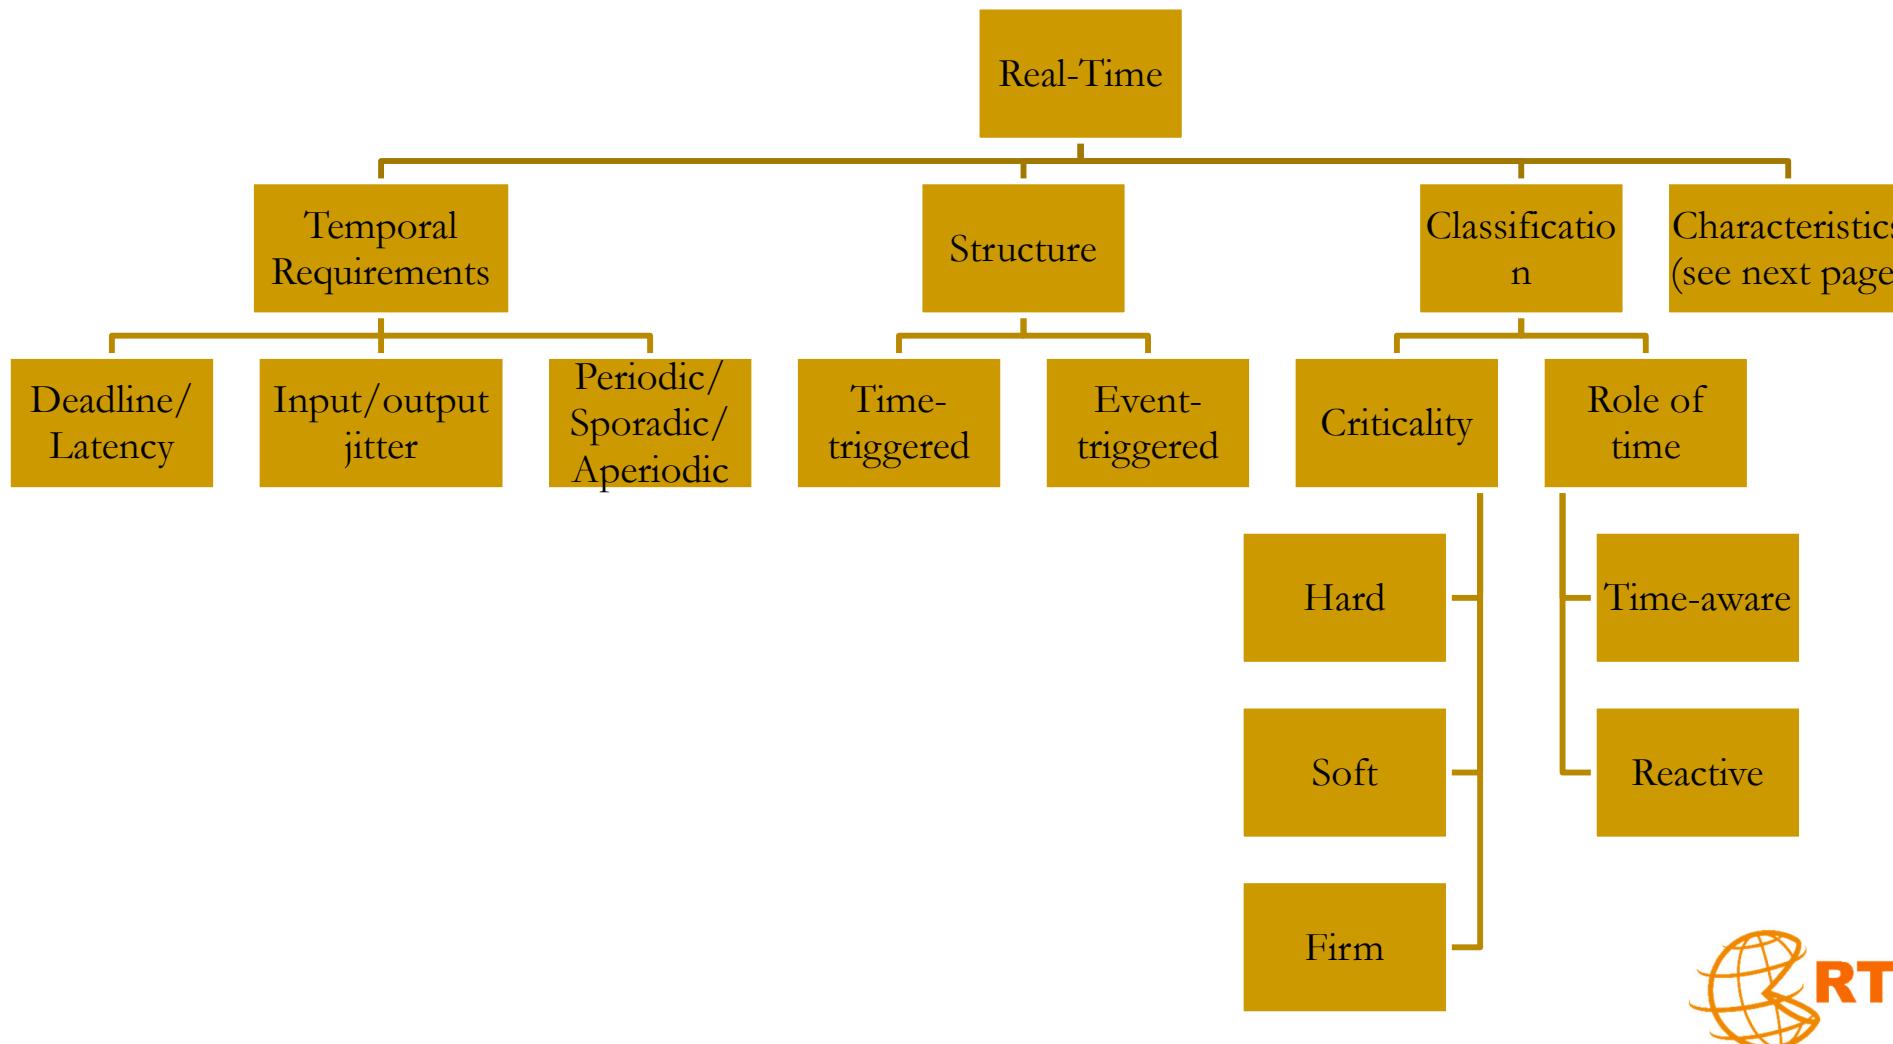
- The design and implementation of an RTS mind the worst case before considering the average case (if at all)
  - Improving the average case is of *no use* and it may even be counterproductive
  - The processor cache is an *average-case* device
  - It operates *adversarially* to the needs of any RTS
- Stability of control prevails over fairness
  - The former concern is selective, the other is general
- When feasibility is proven, starvation is of no consequence
  - The non-critical part of the system may experience starvation



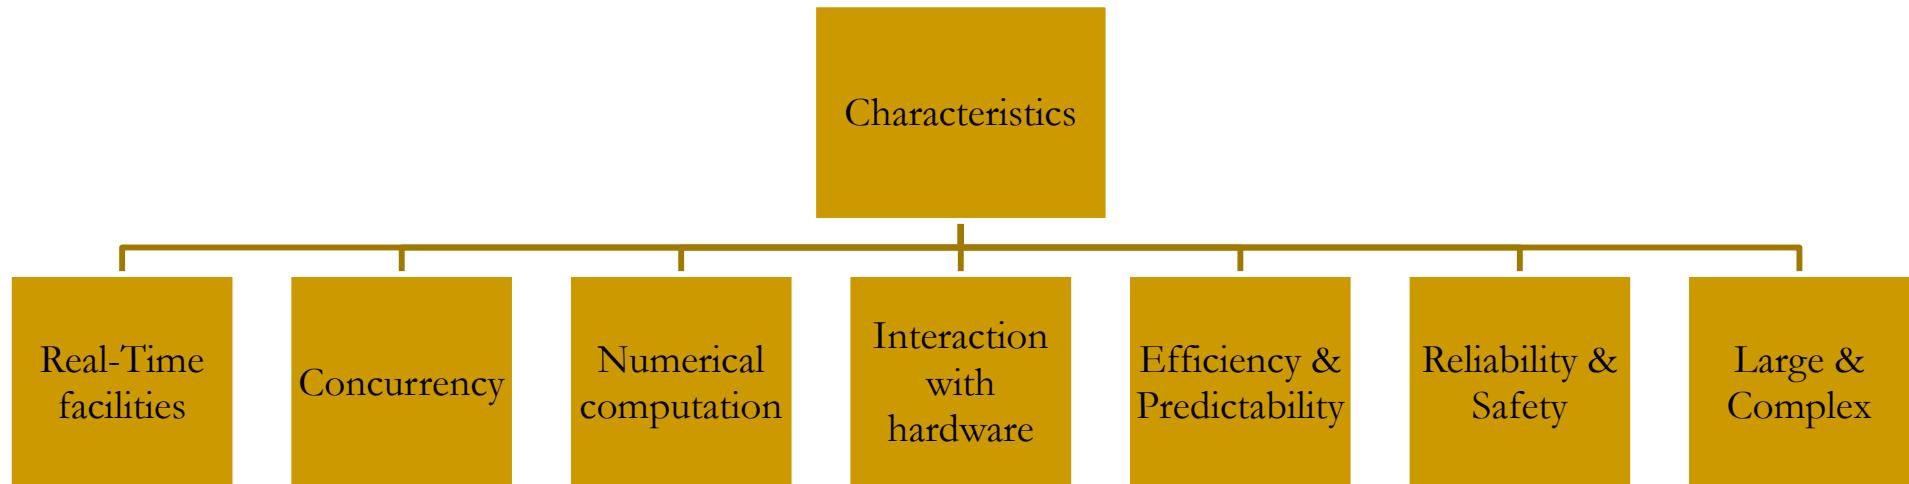
# Summary – 1/3

- From an initial intuition to a more solid definition of real-time embedded systems
- Bird's-eye survey of application requirements and key characteristics
- Taxonomy of tasks
- Abstract models to help reason in general about real-time systems

# Summary – 2/3



# Summary – 3/3



## 2. Scheduling basics

---

**Where we begin to familiarize with the scheduling algorithms that decide how the CPU is assigned to tasks which have a job ready for execution**

# Common approaches – 1/3

## ■ *Clock-driven (time-driven) scheduling*

- Scheduling decisions are made at design time and actuated at fixed time instants during execution by a *scheduler*
  - Such time instants occur at intervals signaled by *clock* via *interrupts*
- The scheduler dispatches to execution (calls) the job due in the current time interval and then suspends itself until the next schedule time
  - The scheduler *is* the prime actor: the jobs are mere called procedures
- Jobs must complete within the assigned time intervals
  - This manner of scheduling does not require *preemption*
  - All scheduling parameters must be known in advance
  - The schedule, computed offline, is fixed forever
  - The scheduling overhead incurred at run time is very small

# Common approaches – 2/3

## ■ *Weighted round-robin scheduling*

- With basic round-robin (which requires preemption)
  - All the ready jobs are placed in a FIFO queue
  - CPU time is quantized, i.e., assigned in slices
  - Job at head of queue is dispatched to execution for one time slice
    - If not complete by end of slice, it goes to tail of queue
    - All jobs in queue are given one single slice per round (full traversal of queue)
  - Not fit for jobs with precedence relations ( $\tau_i$  must complete before  $\tau_j$  may start) whose execution needs multiple slices
  - Fit for producer-consumer pipelines that progress in one-slice increments
- With weighted correction to it (used in network scheduling)
  - Jobs are assigned CPU time according a (fractional) ‘weight’ attribute
  - Job  $J_i$  gets  $\omega_i$  time slices per round
  - One full round corresponds to  $\sum_i \omega_i$  progress for the ready jobs

# Common approaches – 3/3

## ■ *Priority-driven (event-driven) scheduling*

- A *greedy* class of algorithms
  - Never leave available processing resources unused if they are wanted
  - An available resource may stay unused only if no job is ready to use it
  - May not always be good: *clairvoyancy* may show that deferring an assignment of CPU may improve response time (Why?)
  - Anomalies may occur when job parameters change dynamically
- The jobs that contend for execution are kept in a *ready queue*
- Scheduling takes place when the ready queue changes
  - Such scheduling events are called *dispatching points*
  - Scheduling decisions are made online
  - Dispatching employs *preemption*

# Preemption vs. non preemption

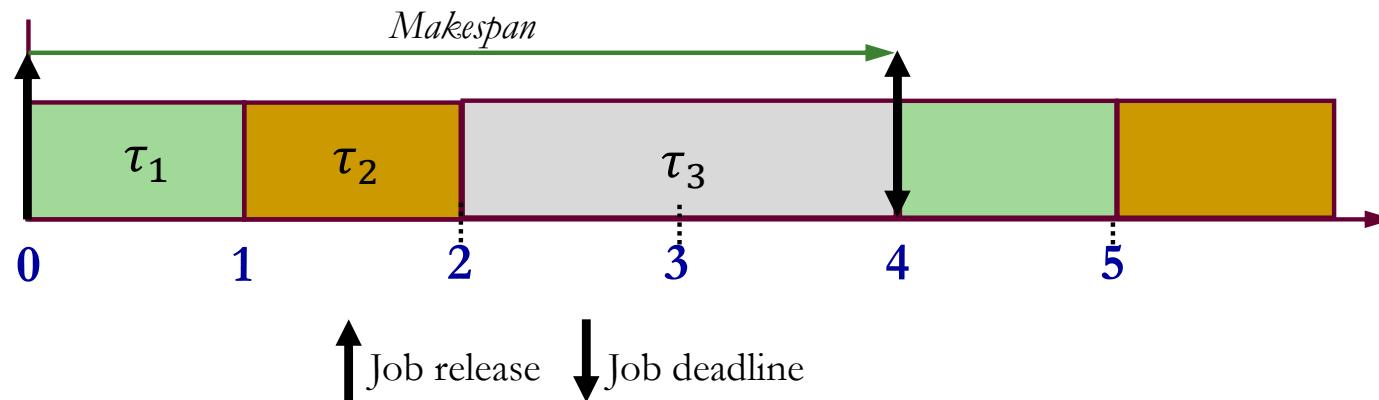
- Can we compare preemptive to non-preemptive scheduling in terms of performance?
  - No single response for all cases
  - Preemptive scheduling *provably better* than non-preemptive when all jobs have same release time and preemption overhead is negligible
    - But preemption costs hardly are negligible
- Does the improvement in the last finishing time (*minimum makespan*) under preemptive scheduling pay off the time overhead of preemption?
  - We do *not* know in general ...
  - For 2 CPUs, the minimum makespan for non-preemptive scheduling is *never worse* than  $4/3$  of that for preemptive

# Makespan

$$\tau_1 = \{p_1 = 4, e_1 = 1, D_1 = p_1\}$$

$$\tau_2 = \{p_2 = 4, e_2 = 1, D_2 = p_2\}$$

$$\tau_3 = \{p_3 = 4, e_3 = 2, D_3 = p_3\}$$

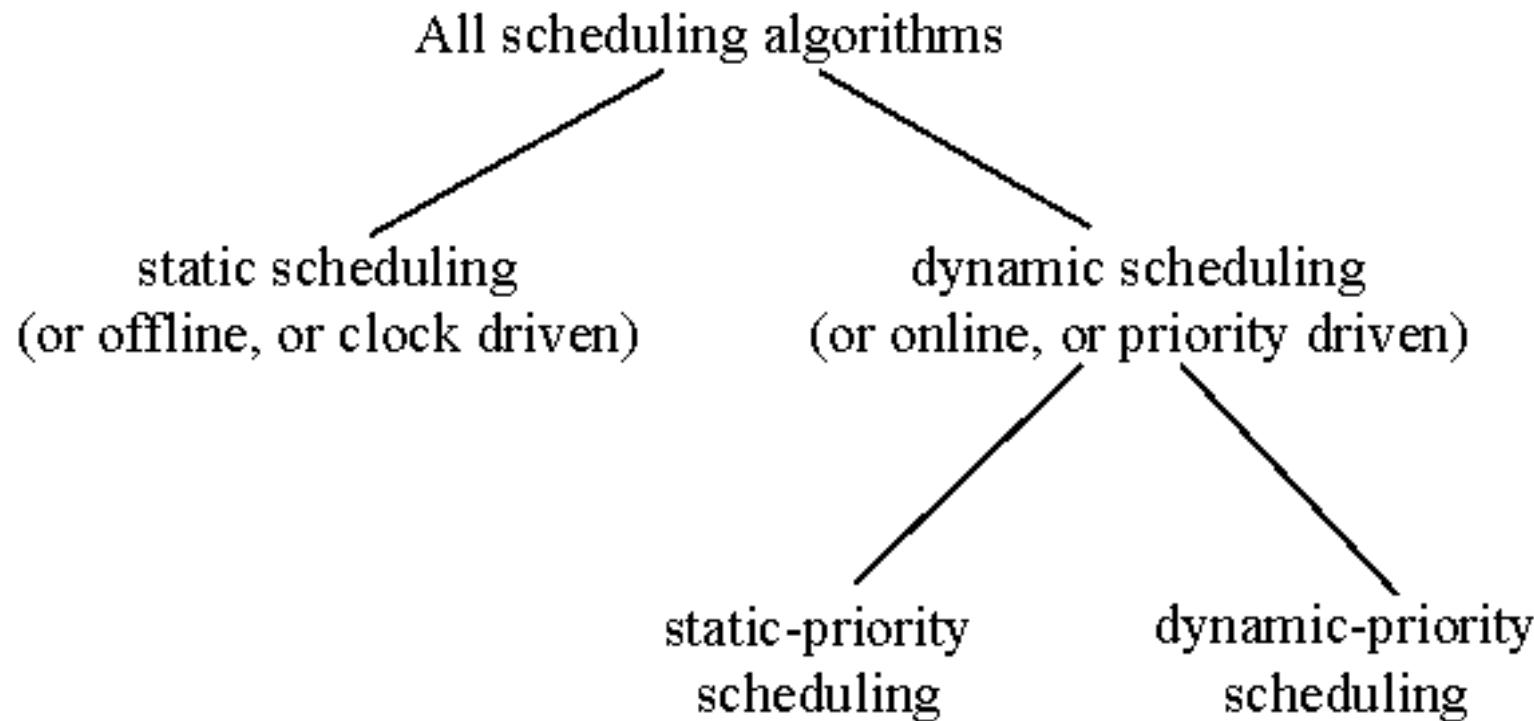


From the release of the first job of the first task to the completion of the first job of the last task

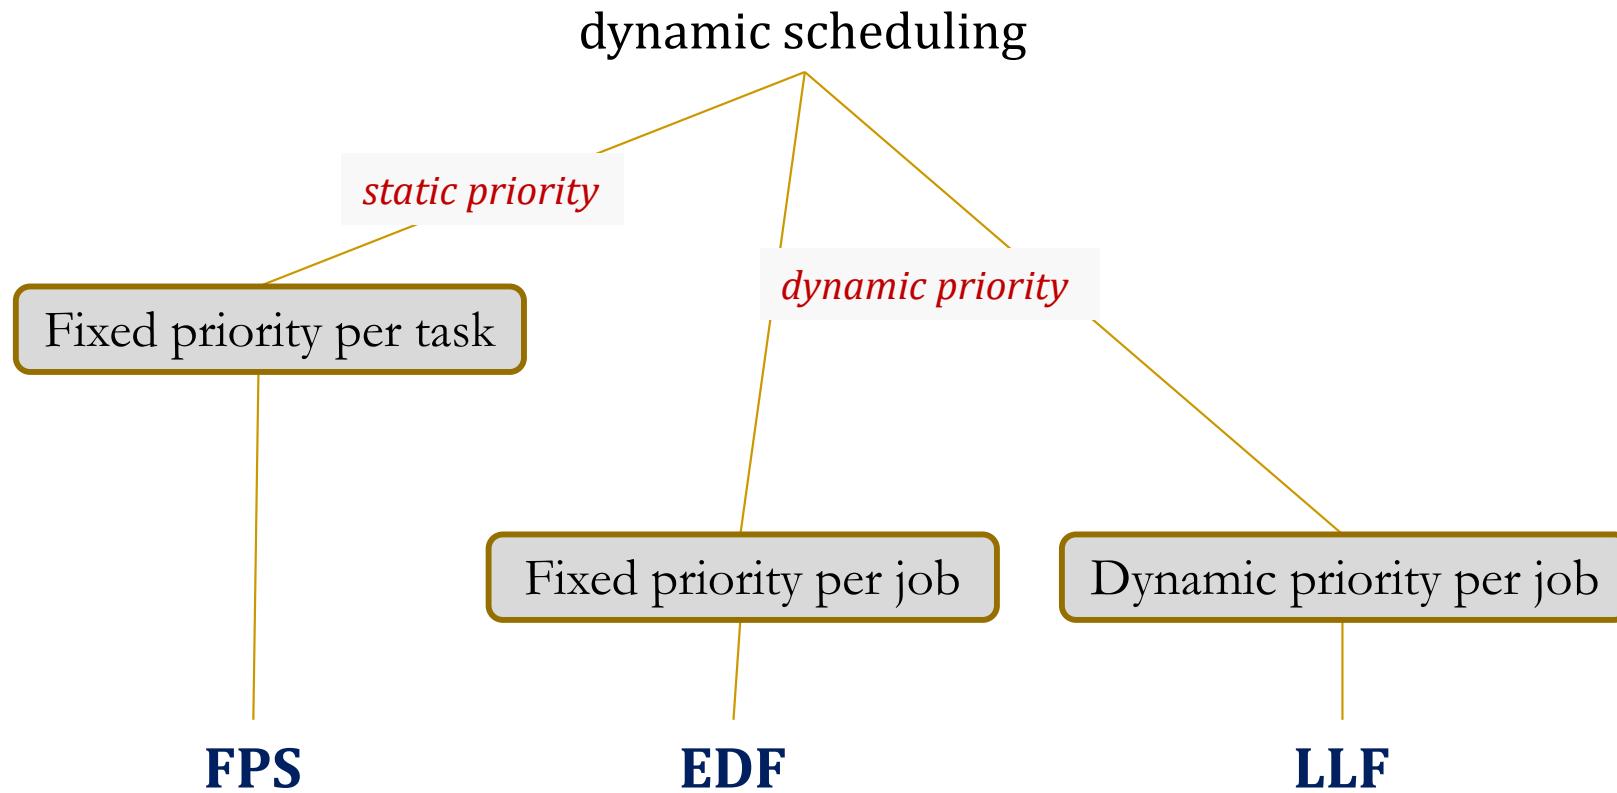
# Interpreting predictability of execution

- The execution of system  $S$  under a given scheduling algorithm  $A$  is *predictable* if the start time and response time of every job in  $S$  strictly vary between a *minimal* and a *maximal schedule*
  - *Maximal schedule* is created by  $A$  under worst-case conditions for contention and execution demands
  - *Minimal schedule*, analogously for the best case
- Predictable scheduling algorithms can be reasoned about in terms of worst-case behavior!
- **Theorem:** the execution of *independent* jobs with given release times under preemptive priority-driven scheduling on a single processor **is predictable**

# Classification of Scheduling Algorithms

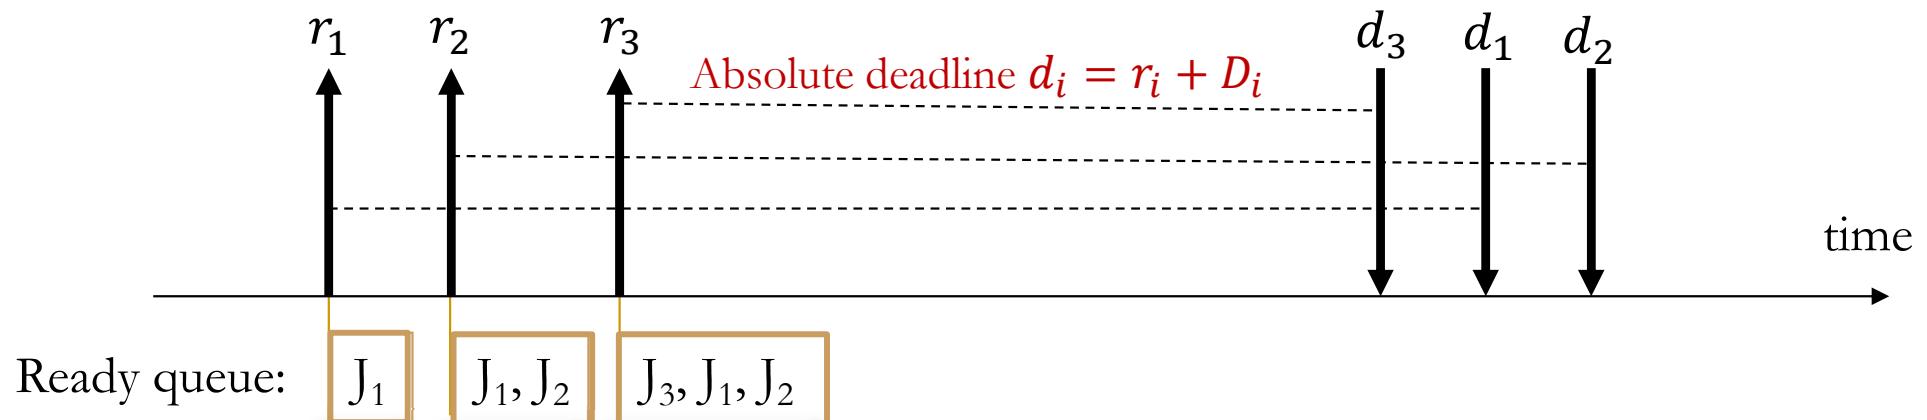


# Taxonomy of dynamic scheduling



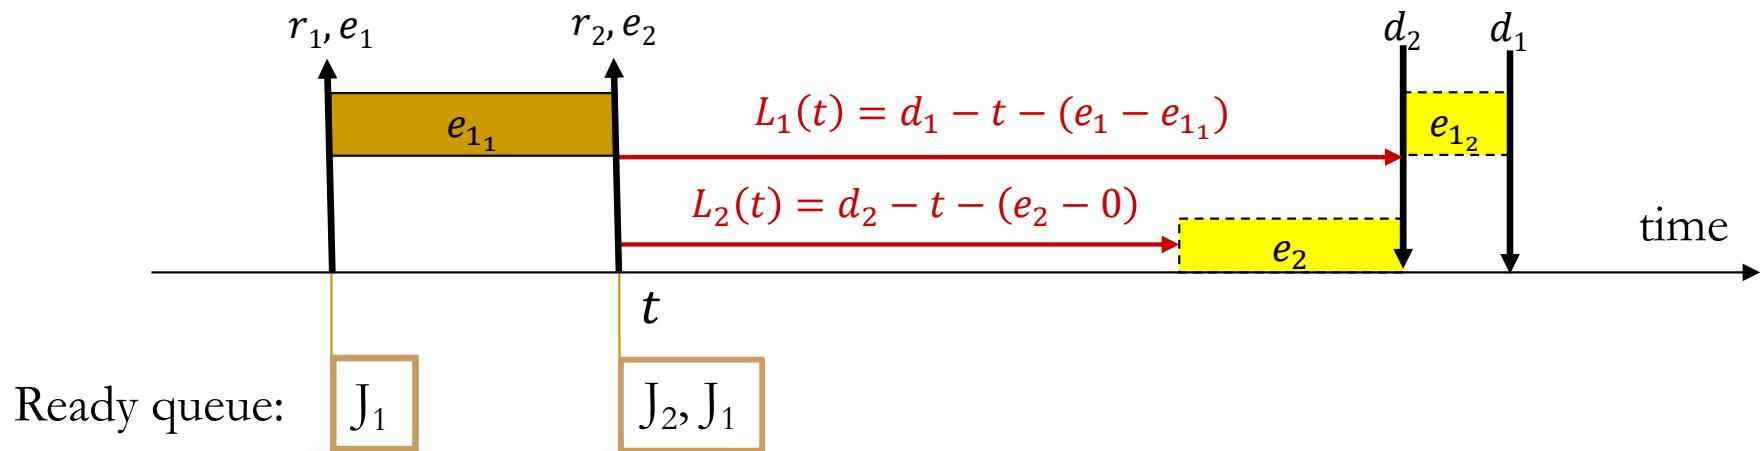
# Ways to optimality – 1/2

- **Earliest Deadline First** (EDF) scheduling assigns priorities in inverse relation to *absolute deadlines*
  - This makes it dynamic per task but static per job
  - Ready queue changes on job release and on job completion
- [Liu & Layland: 1973]: EDF scheduling is *optimal* for single-CPU systems with independent jobs and preemption
  - For any job set, EDF produces a feasible schedule if one exists
  - The optimality of EDF breaks otherwise (e.g., no preemption, parallelism)



# Ways to optimality – 2/2

- **Least Laxity First** (LLF) scheduling assigns priorities in inverse relation to *laxity*  $L(t)$ 
  - Jobs' priority,  $L(t)$ , varies with  $t$ : dynamic per job and more costly than EDF
  - $L_i(t) = d_i - t - Y_i(t)$ , where  $Y_i(t)$  is the residual execution time needed for  $\tau_i$  at time  $t$ , with release time  $r_i$  and relative deadline  $D_i$  ( $d_i = r_i + D_i$ )
  - Ready queue reordering occurs on job release and job completion
- [Liu & Layland: 1973]: LLF is as *optimal* as EDF



# Clock-driven scheduling – 1/8

## ■ *Workload model*

- $N$  periodic tasks, for  $N$  constant and statically defined
- The  $(\varphi_i, p_i, e_i, D_i)$  parameters of every task  $\tau_i$  are constant and statically known
- The schedule is static and committed at design to a table  $S$  of **decision times**  $t_k$ , where:
  - $S[t_k] = \tau_i$  if a job of task  $\tau_i$  must be dispatched at time  $t_k$
  - $S[t_k] = I$  (*idle*) if no job is due at time  $t_k$
  - Schedule computation can be as sophisticated as we like since we pay for it only at design time
  - Jobs *cannot overrun* otherwise the system is in error

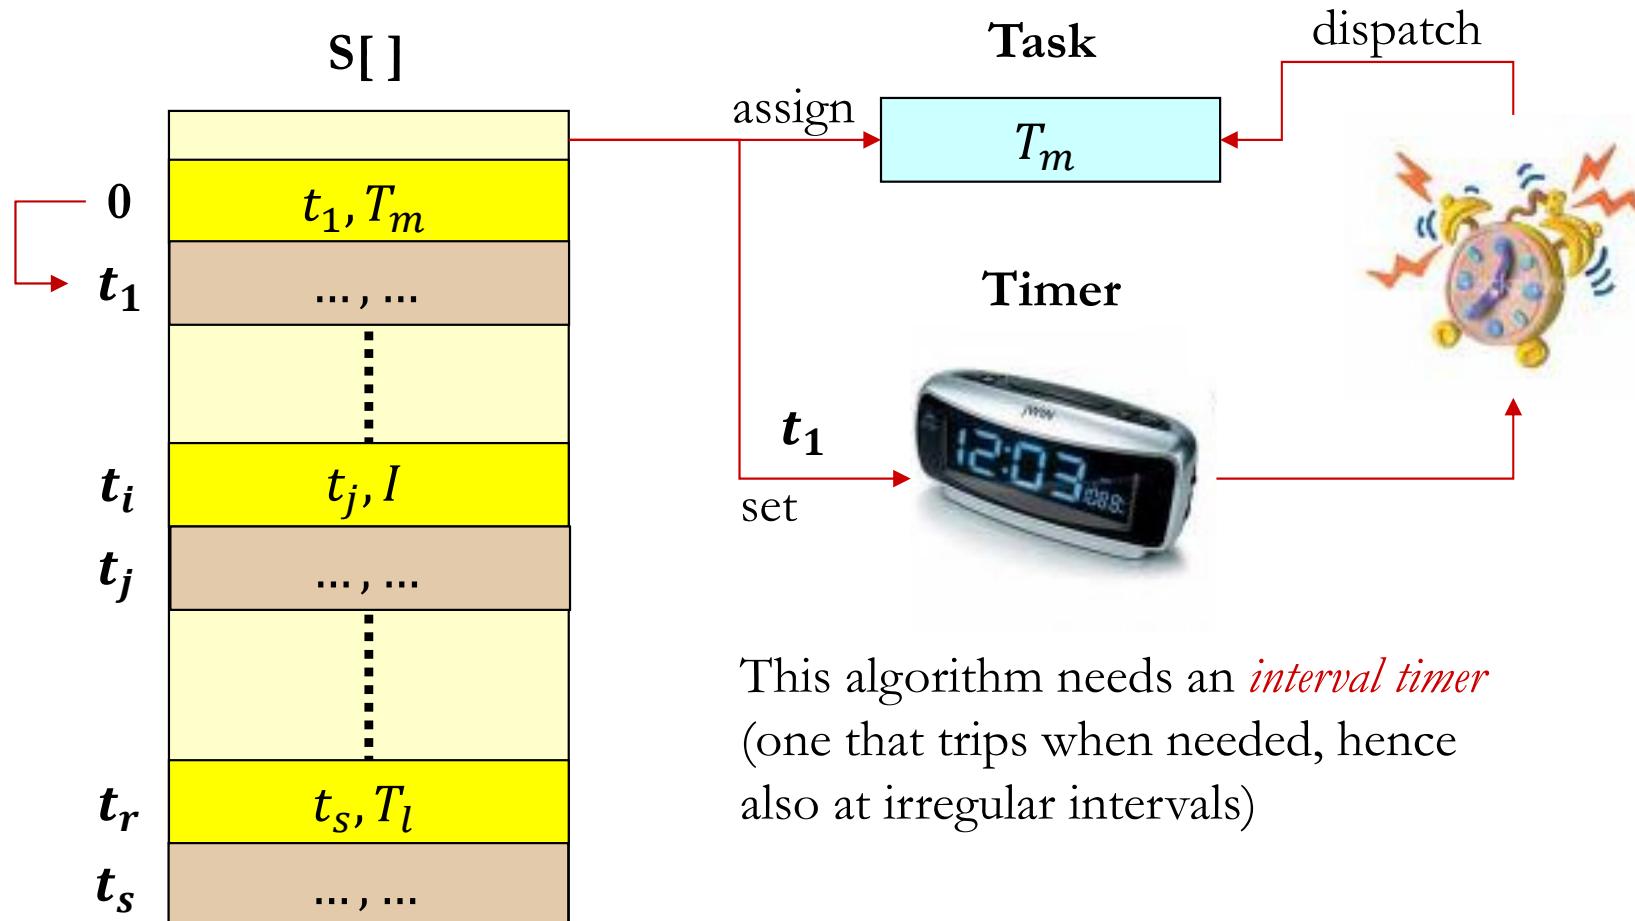
# Clock-driven scheduling – 2/8

**Input:** stored schedule  $S[t_k]$ ,  $k = \{0, \dots, N - 1\}$ ;  $H$  (hyperperiod)

**SCHEDULER ::**

```
i := 0; // slot number
k := 0; // N-modulo correspondent if index i
set timer to expire at  $t_k$  ;
do forever :
    sleep until timer interrupt;
    if aperiodic job executing then preempt; end if;
    T :=  $S[t_k]$  ; // current task
    i := i + 1;
    k := i mod N;
    set timer to expire at  $t_k + \lfloor i/N \rfloor \times H$ ;
    if  $T = I$  then execute job at head of aperiodic queue;
        else execute job of task  $T$ ;
    end if;
end do;
end SCHEDULER
```

# Clock-driven scheduling – 3/8



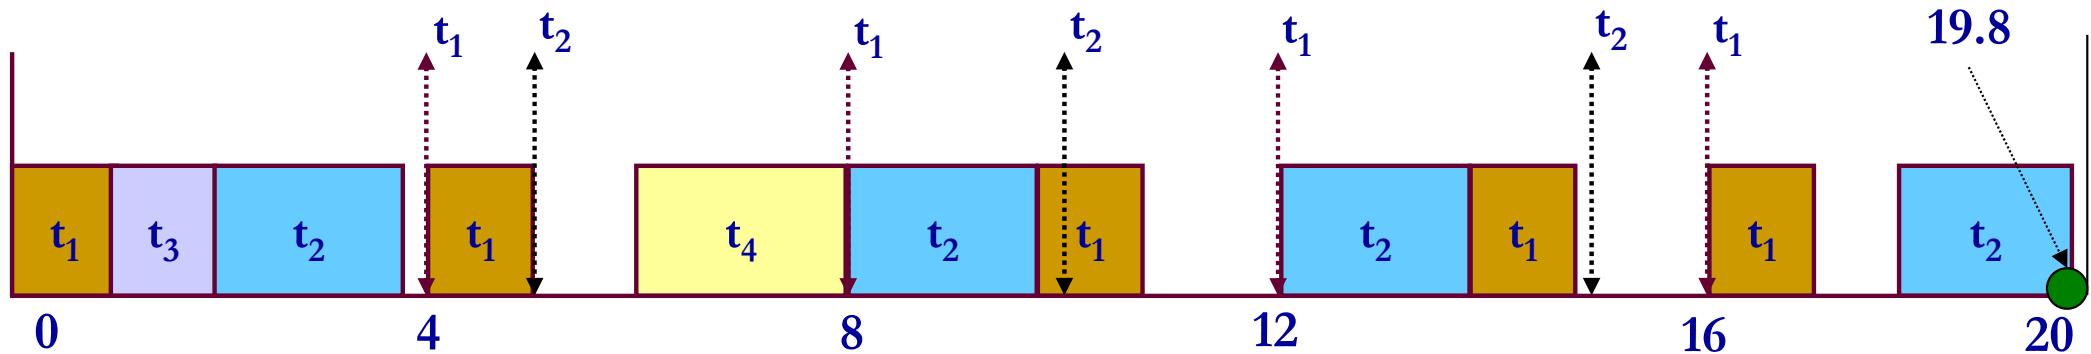
Where the  $t_j$  values need *not* be equally spaced

# Example

$$(\varphi_i, p_i, e_i, D_i)$$

$$J = \{t_1 = (0, 4, 1, 4), t_2 = (0, 5, 1.8, 5), t_3 = (0, 20, 1, 20), t_4 = (0, 20, 2, 20)\}$$

$$U = \sum_i \frac{e_i}{p_i} = 0.76, H = 20$$



Time	Schedule
0	$t_1$
1	$t_3$
2	$t_2$
3.8	I
4	$t_1$
...	...
19.8	I
20	Goto $t \bmod(H)$

- The schedule table S for J would need 17 entries
  - That's too many and the schedule too fragmented!
- Why 17?

# Clock-driven scheduling – 4/8

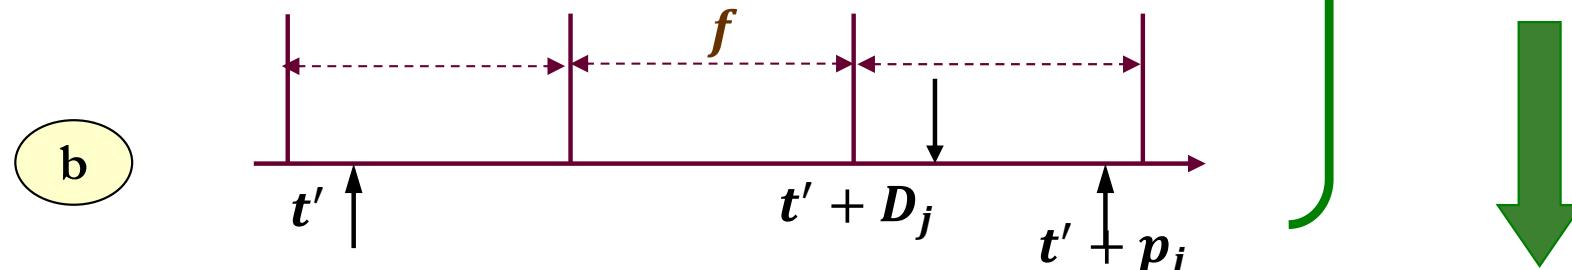
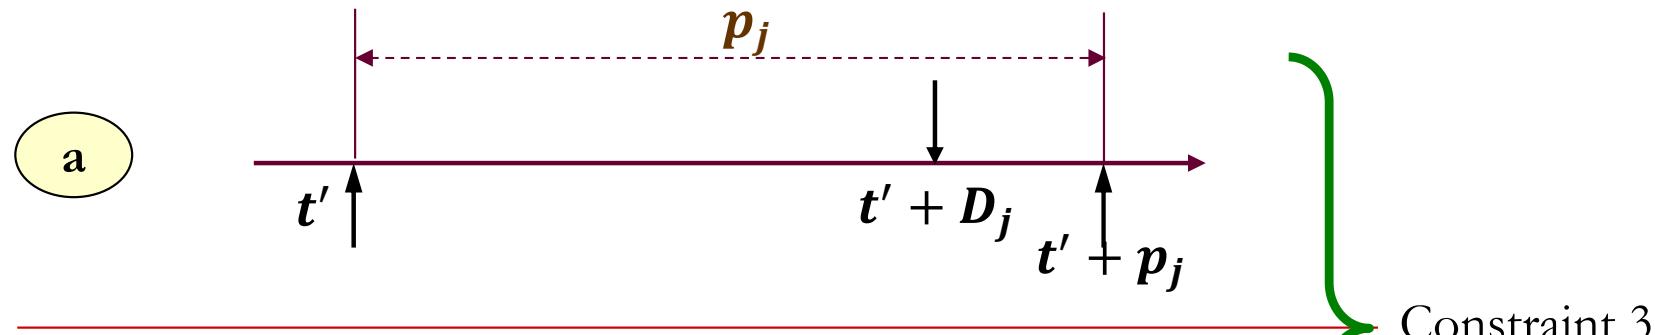
- To keep design complexity low, we should *minimize* the size of the cyclic schedule (table  $S$ ): how can that be done?
- The scheduling point  $t_k$  should occur at regular intervals
  - Each such interval is termed ***minor cycle*** (*frame*) and has duration  $f$
  - We need a (cheaper, more standard) ***periodic timer*** instead of a (more costly) interval timer
  - Within minor cycles there is no preemption, but a single frame may allow the execution of multiple (run-to-completion) jobs
- For every task  $\tau_i$ ,  $\varphi_i$  must be a non-negative integer multiple of  $f$ 
  - Forcibly, the first job of every task has its release time set at the start edge of a minor cycle
- To build such a schedule, we must enforce some constraints

# Clock-driven scheduling – 5/8

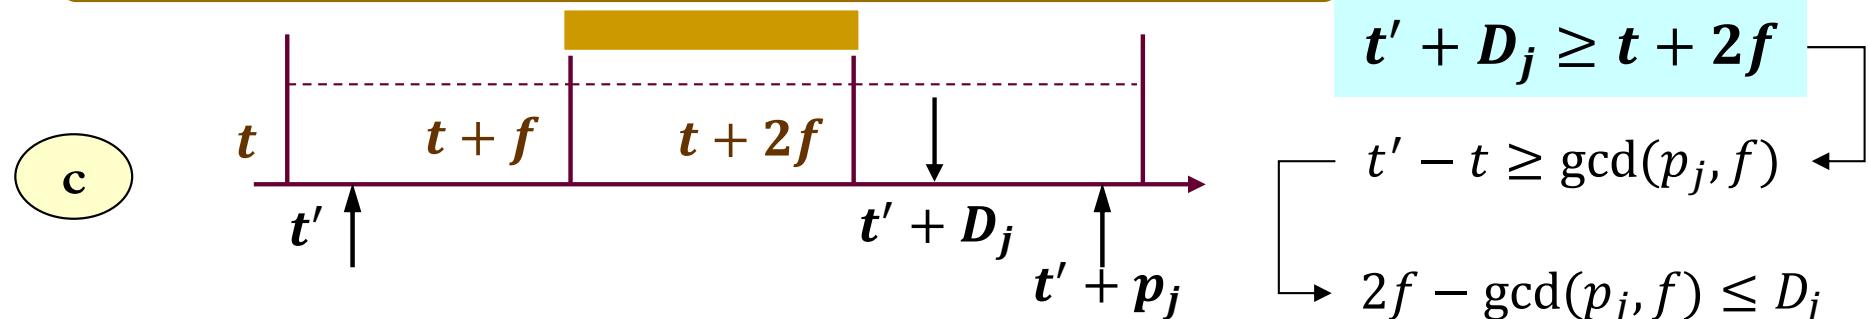
- **Constraint 1:** Every job  $J$  must complete within  $f$ 
  - $f \geq \max_{i=\{1,\dots,n\}}(e_i)$  so that *overruns* can be detected
- **Constraint 2:**  $f$  must be integer divisor of hyperperiod
  - $H : H = Nf$  where  $N \in \mathbb{N}$
  - It suffices that  $f$  be an integer divisor of at least one task period  $p_i$
  - The hyperperiod beginning at minor cycle  $kf$  for  $k = 0, N - 1, 2N - 1$  is termed *major cycle*
- **Constraint 3:** There must be one *full* frame  $f$  between  $J$ 's release time  $t'$  and its deadline:  $t' + D_j \geq t + 2f$ 
  - So that  $J$  can be scheduled and complete in that one frame
  - This can be expressed as:  $2f - \gcd(p_i, f) \leq D_i$  for every task  $\tau_i$

# Understanding constraint 3

Finding the worst *allowable* offset between  $t'$  (the job's release time) and  $\mathbf{t}$  (the edge of the frame)



This is the frame in which job  $J$  must be scheduled



# Example

$$(\varphi_i, p_i, e_i, D_i)$$

- $T = \{(0, 4, 1, 4), (0, 5, 2, 5), (0, 20, 2, 20)\}$
- $H = 20$
- [c1] :  $f \geq \max(e_i) : f \geq 2$
- [c2] :  $\lfloor p_i/f \rfloor - p_i/f = 0 : f = \{2, 4, 5, 10, 20\}$
- [c3] :  $2f - \gcd(p_i, f) \leq D_i : f \leq 2$

$$f = 2 : 4 - \gcd(4, 2) \leq 4 \text{ OK}$$

$$4 - \gcd(5, 2) \leq 5 \text{ OK}$$

$$4 - \gcd(20, 2) \leq 20 \text{ OK}$$

$$f = 4 : 8 - \gcd(4, 4) \leq 4 \text{ OK}$$

$$8 - \gcd(5, 4) \leq 5 \text{ KO}$$

$$f = 5 : 10 - \gcd(4, 2) \leq 4 \text{ KO}$$

$$f = 10 : 20 - \gcd(4, 2) \leq 4 \text{ KO}$$

$$f = 20 : 40 - \gcd(4, 2) \leq 4 \text{ KO}$$

# Clock-driven scheduling – 6/8

- It is very likely that the original parameters of some task set  $T$  may prove unable to satisfy all three constraints for any given  $f$  simultaneously
- In that case, we must decompose the jobs of for some task(s)  $\tau_i$  by **slicing** their (WCET)  $e_i^W$  into fragments artificially sized to yield a “good”  $f$

# Clock-driven scheduling – 7/8

- To construct a cyclic schedule, we must make three design decisions
  - Fix an  $f$
  - Slice (the large) jobs
  - Assign (jobs and) slices to minor cycles
- Sadly, these decisions are very tightly coupled
  - This defect makes cyclic scheduling *very* fragile to any change in system parameters

# Clock-driven scheduling – 8/8

**Input:** stored schedule  $S[k]$ ,  $k$  in  $0 \dots F - 1$

**CYCLIC\_EXECUTIVE ::**

$t := 0; k := 0;$

**do forever**

**sleep until** clock interrupt at time  $t \times f$ ;

    currentBlock :=  $S[k]$ ;

$t := t + 1; k := t \bmod F$ ;

**if** last job not completed **then** take action;

**end if**;

    execute all job slices in currentBlock;

**while** aperiodic job queue not empty **do**

        execute aperiodic job at top of queue;

**end do**; // the clock interrupt may occur!

**end do**;

**end SCEDULER**

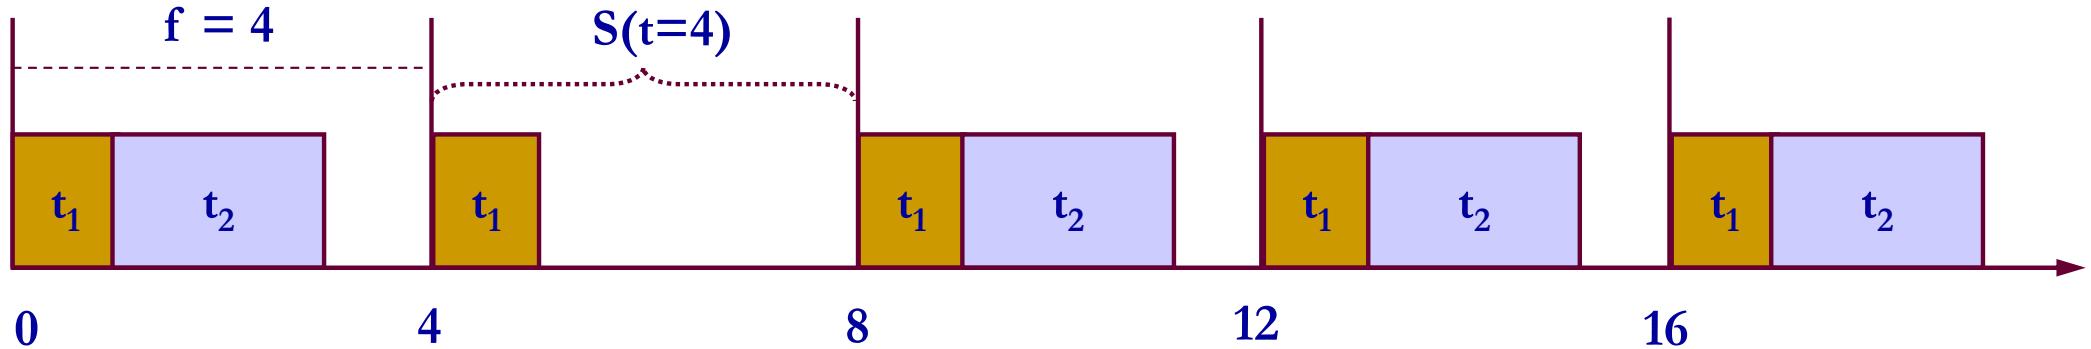
# Example (slicing) – 1/2

$$(\varphi_i, p_i, e_i, D_i)$$

$$J = \{\tau_1 = (0, 4, 1, 4), \tau_2 = (0, 5, 2, 5), \tau_3 = (0, 20, 5, 20)\}, H = 20$$

$\tau_3$  causes disruption since we need  $e_3 \leq f \leq 4$  to satisfy c3

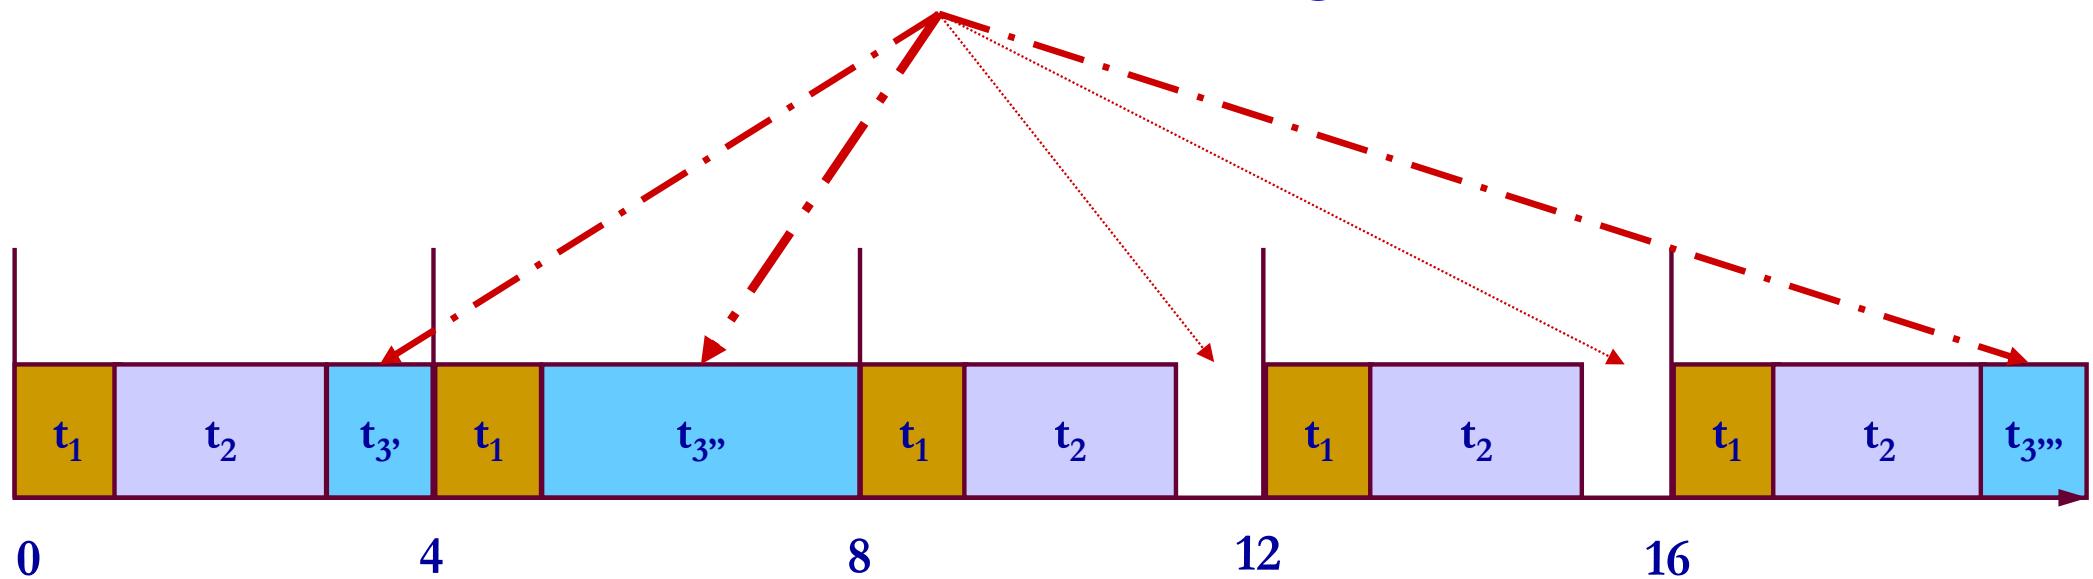
We must therefore slice  $e_3$  : how many slices do we need?



We first look at the schedule with  $f = 4$  and  $F = \left(\frac{H}{f}\right) = 5$  without  $\tau_3$ , to see what least-disruptive opportunities we have ...

# Example (slicing) – 2/2

... then we observe that  $e_3 = \{1, 3, 1\}$  is a good choice



$$\tau_3 = \{\tau'_3 = (0, 20, 1, x), \tau''_3 = (0, 20, 3, y), \tau'''_3 = (0, 20, 1, 20)\}$$

where  $x < y \leq 20$  represent the precedence constraints that must hold between the slices (could have used phases instead)

# Design issues – 1/3

- Completing a job much ahead of its deadline may be of little or no use!
- Spare time in time slices should be given to *aperiodic jobs*, to allow the system to produce more (piece-rate) value added
- This can be achieved with ***slack stealing***, which allows aperiodic jobs to execute *in preference* to periodic jobs when possible
  - Minor cycles may have *slack* time, not used for scheduling periodic jobs, that may be devolved to aperiodic jobs
  - The slack of minor cycles is a *statically* known at design time
- A cyclic scheduler does slack stealing if it assigns the available slack time *at the beginning* of every minor cycle (instead of at the end)
  - The system becomes *more reactive* (which is good)
  - A fine-grained interval timer signals end of slack time for each minor cycle

# Design issues – 2/3

- What can we do to handle ***overruns***?
  - Halt the job found running at the start of the new minor cycle
    - But that job may not be the one that overrun!
    - Even if it was, stopping it would only serve a useful purpose if producing a late result had no residual *utility*
  - Defer halting until the job has completed all its “critical actions”
    - To avoid the risk that a premature halt may leave the system in an inconsistent state
  - Allow the job some extra time by delaying the start of the next minor cycle
    - Plausible if producing a late result still had *utility*

# Design issues – 3/3

- What can we do to handle ***mode changes***?
  - A mode change is when the system incurs some reconfiguration of its function and workload parameters
- Two main axes of design decisions
  - With or without deadline during the transition
  - With or without overlap between outgoing and incoming operation modes

# Overall evaluation

## ■ Pros

- Comparatively simple design
- Simple and robust implementation
- Complete and cost-effective verification

## ■ Cons

- Very fragile design
  - Construction of the schedule table is a NP-hard problem
  - High extent of undesirable architectural coupling
- All parameters must be fixed *a priori* at the start of design
  - Choices may be made arbitrarily to satisfy the constraints on  $f$
  - Totally inapt for sporadic jobs

# Priority-driven scheduling

- Base principle
  - Every job is assigned a priority
  - The job with the highest priority is dispatched to execution
- Two implementation decisions
  - When jobs' priority should change
  - When dispatching should occur
- ***Dynamic-priority scheduling***
  - Distinct jobs of the same task may have *distinct* priorities
    - EDF: the job priority is *fixed* at release, but changes across releases
    - LLF: the job priority may change at every dispatching point
- ***Static-priority scheduling***
  - All jobs of the same task have *one and the same* priority

# Fixed-priority scheduling (FPS)

- Two main strategies for priority assignment
  - This is all we need to determine FPS
- ***Rate monotonic*** (when deadlines are implicit)
  - A task with *faster rate* (lower period) takes precedence
  - Optimal *static-per-task* priority assignment under preemptive scheduling and implicit deadlines
  - The consequent scheduling is called **RMS**
- ***Deadline monotonic*** (when deadlines are constrained)
  - A task with *higher urgency* (shorter relative deadline) goes first
  - Optimal (static-per-task) for constrained deadlines
- In effect, both algorithms are deadline-based ☺

# Preliminary observations

- Priority-driven scheduling algorithms that disregard job urgency (deadline) perform *poorly*
- The WCET is *not* a factor of consequence for priority assignment
  - Weighed round-robin scheduling is “utilization-monotonic”, but is unfit for real-time systems
- **Schedulable utilization** is a good metric to compare the performance of scheduling algorithms
  - A scheduling algorithm  $S$  can produce a feasible schedule for a task set  $T$  on a single processor if and only if  $U(T)$  does not exceed the schedulable utilization of  $S$

# Utilization-based tests – 1/2

- **Theorem** [Liu & Layland: 1973]

For single processors and implicit deadlines, EDF's *schedulable utilization*  $U(n) = \sum_{i=1}^n \frac{e_i}{p_i} \leq 1$  is *always* feasible

- A *necessary and sufficient* (exact) test condition for EDF

- For RMS, the equivalent test is

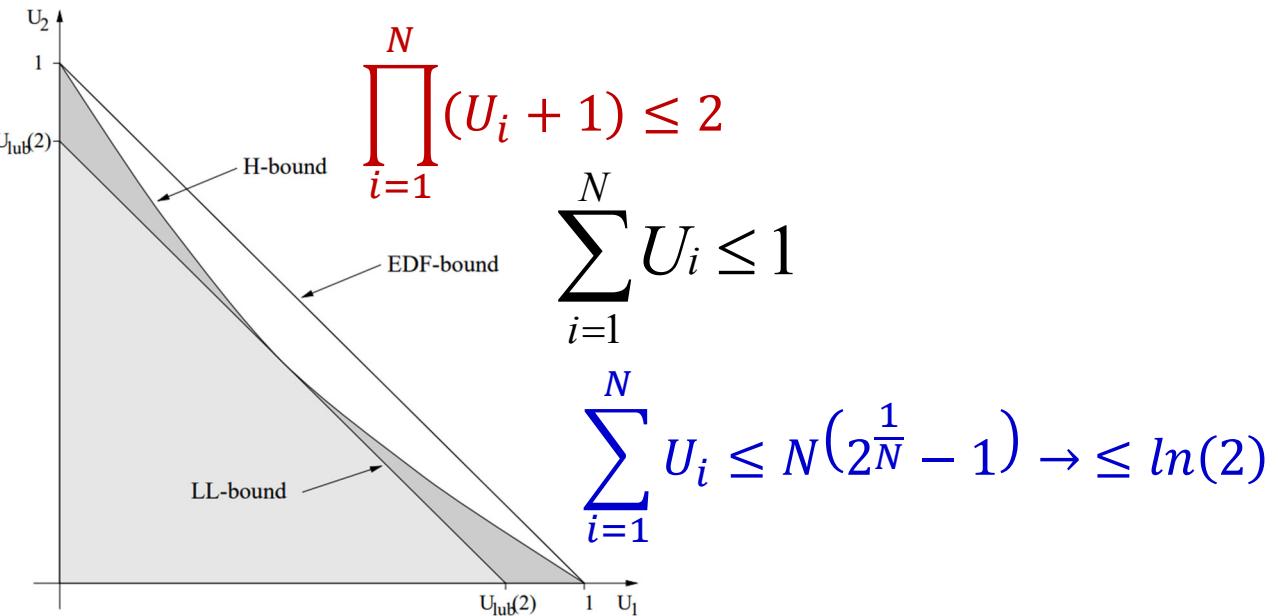
$$U(n) = \sum_{i=1}^n \frac{e_i}{T_i} \leq n\left(2^{\frac{1}{n}} - 1\right)$$
$$\lim_{n \rightarrow \infty} n\left(2^{\frac{1}{n}} - 1\right) = \ln 2 \sim 0.69$$

- A *sufficient but not necessary* (pessimistic) test for RMS
- The schedulable utilization of RMS is *less* than for EDF

# Utilization-based tests – 2/2

- The *hyperbolic bound* [Bini & Buttazzo, 2001] improves the Liu & Layland utilization test for RMS
  - It helps prove that RMS achieves 100% utilization when *all pairs* of periods in the task set are in harmonic relation

**Examples of feasibility regions**  
Plot in an  $n = 2$  U-space, where each point  $U = \{U_1, U_2, \dots, U_n\}$  represents a periodic task set with utilization  $U_i$

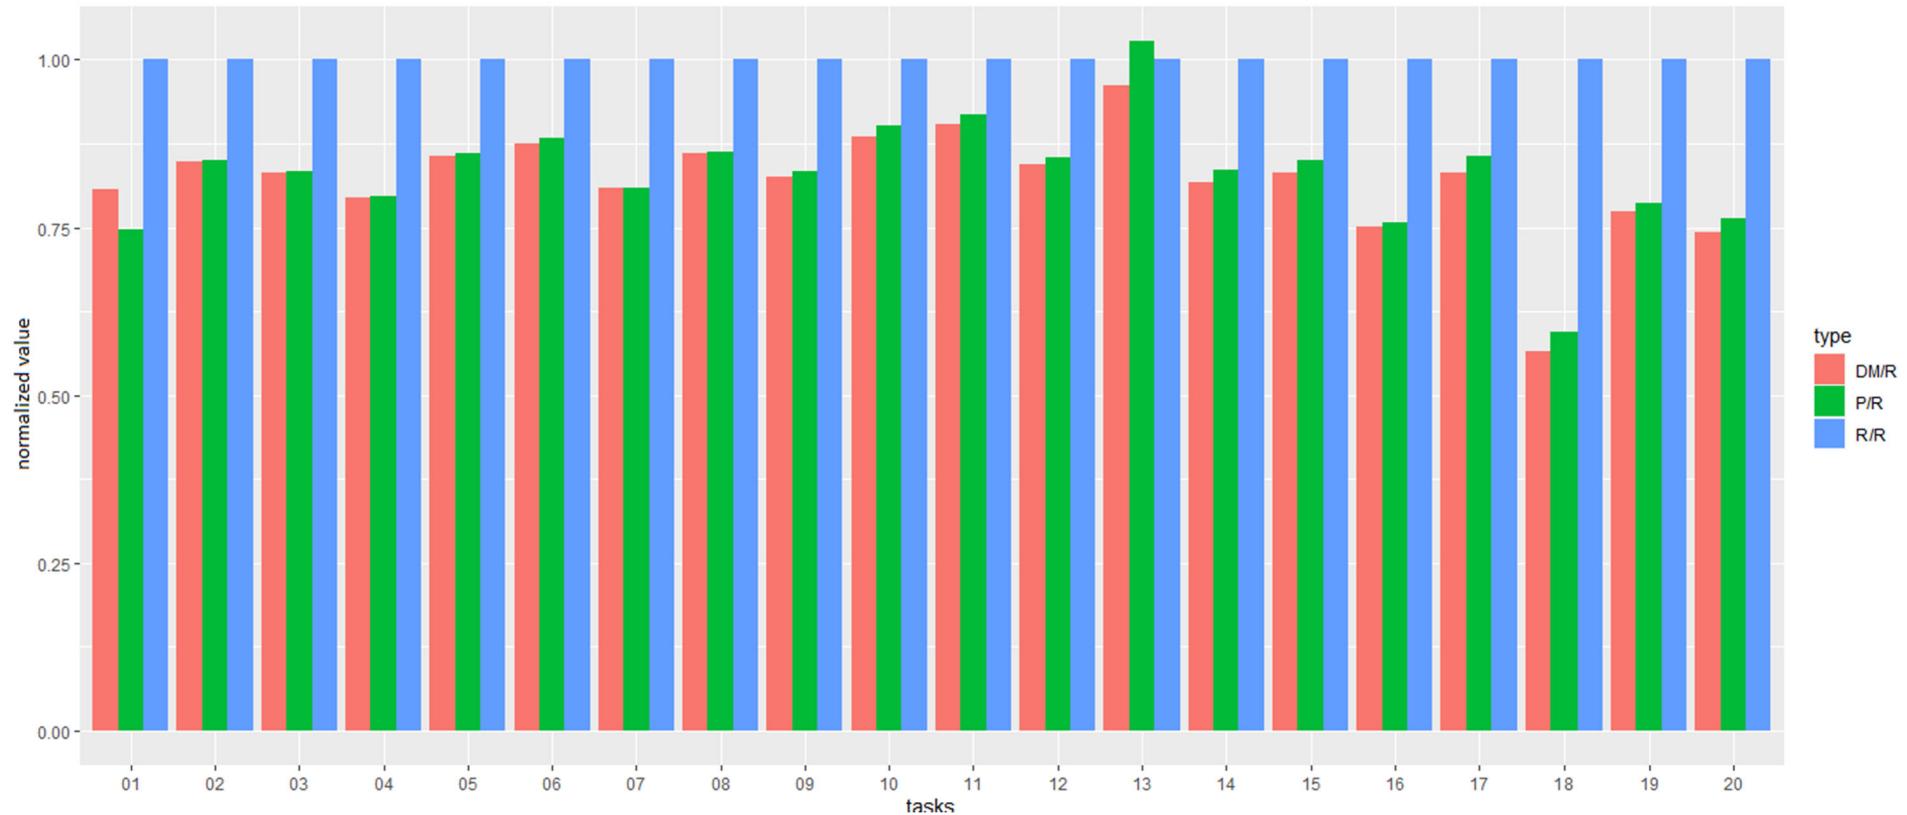


# Appraising feasibility

- Schedulable utilization alone is *not* a sufficient criterion: we must also consider *predictability*
  - Recall its intuition, given in Section 1
- On ***transient overload***, the behavior of static-priority scheduling can be determined a-priori and is reasonable
  - The overrun of any job of a given task  $\tau$  does not harm the tasks with higher priority than  $\tau$
- Under transient overload, EDF acts “strange” (locally)
  - A job that missed its deadline is *more urgent* than a job with a deadline in the future: one lateness may cause many more!

# Overload situations – 1/5

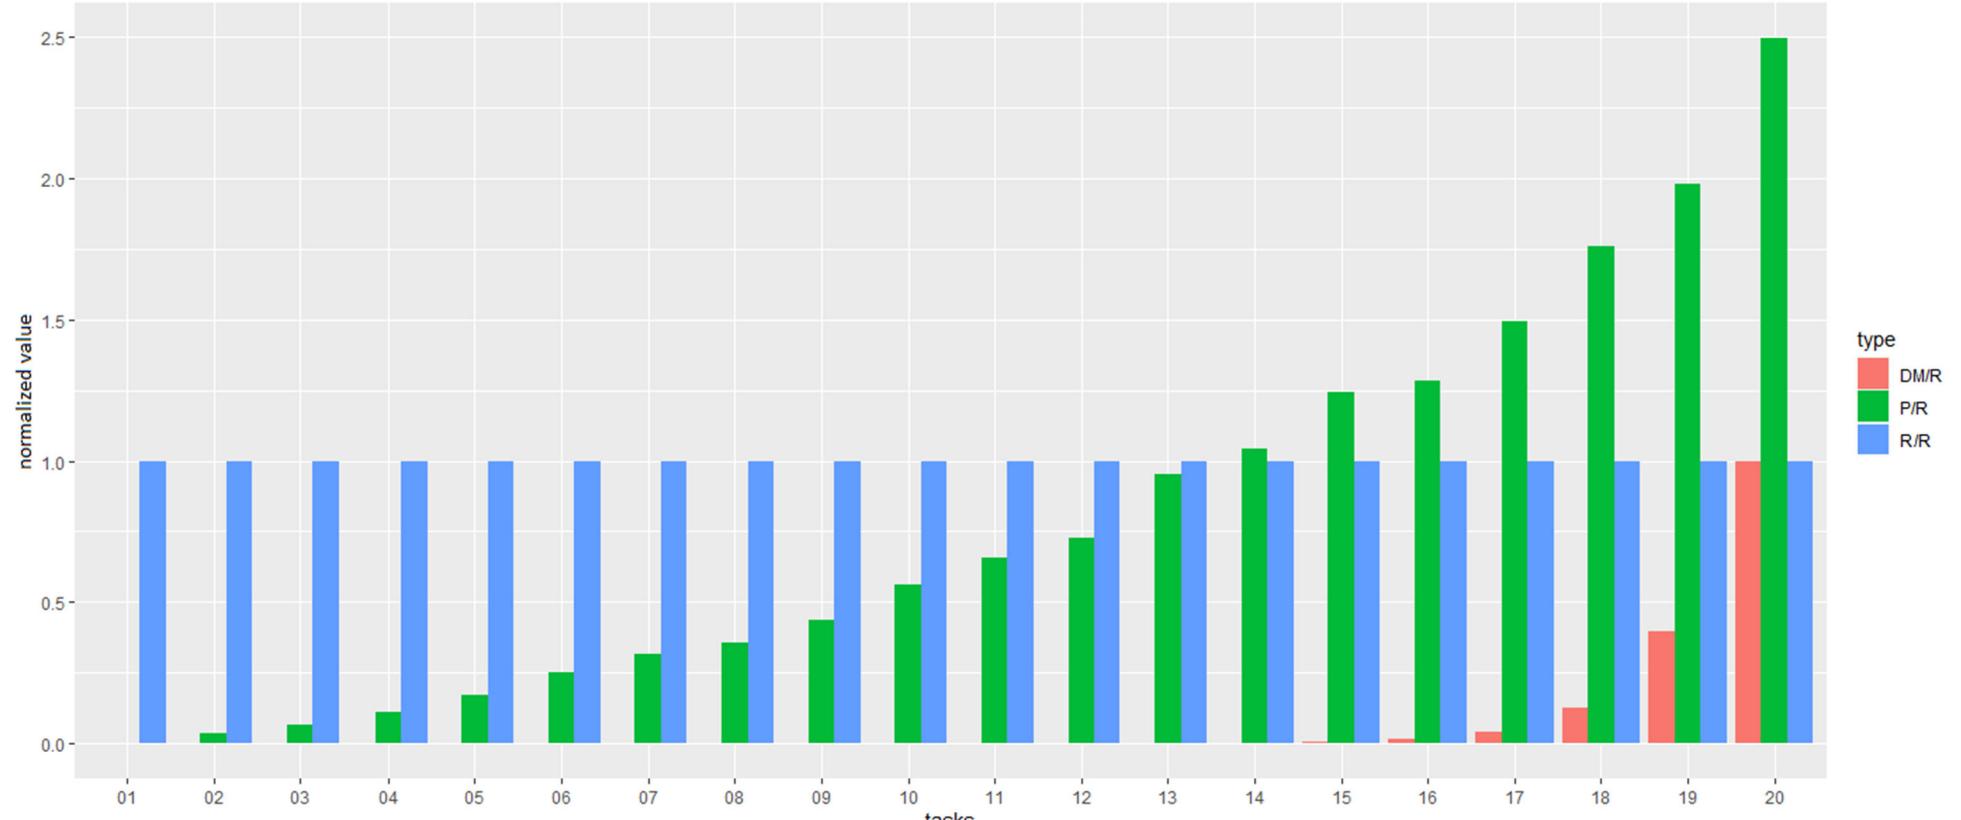
Deadline miss and preemption count ratio over normalized run count (EDF,  $U > 1$ )



Legend: DM/R (deadline misses over releases); P/R (preemptions over releases); R (release; run)

# Overload situations – 2/5

Deadline miss and preemption count ratio over normalized run count (FPS,  $U > 1$ )



Legend: DM/R (deadline misses over releases); P/R (preemptions over releases); R (release; run)

# Overload situations – 3/5

An interesting property of EDF during permanent overloads is that it automatically performs a period rescaling, and tasks start behaving as they were executing at a lower rate. This property has been proved by Cervin et al. (2002) and it is formally stated in the following theorem.

**Theorem 1** [Cervin]. *Assume a set of  $n$  periodic tasks, where each task is described by a fixed period  $T_i$ , a fixed execution time  $C_i$ , a relative deadline  $D_i$ , and a release offset  $\Phi_i$ . If  $U > 1$  and tasks are scheduled by EDF, then, in stationarity, the average period  $\bar{T}_i$  of each task  $\tau_i$  is given by  $\bar{T}_i = T_i U$ .*



Real-Time Systems, 29, 5–26, 2005

© 2005 Springer Science + Business Media, Inc. Manufactured in The Netherlands.

- EDF's throughput decreases by period rescaling
- FPS's throughput decreases by discarding lower-priority jobs

# Overload situations – 4/5

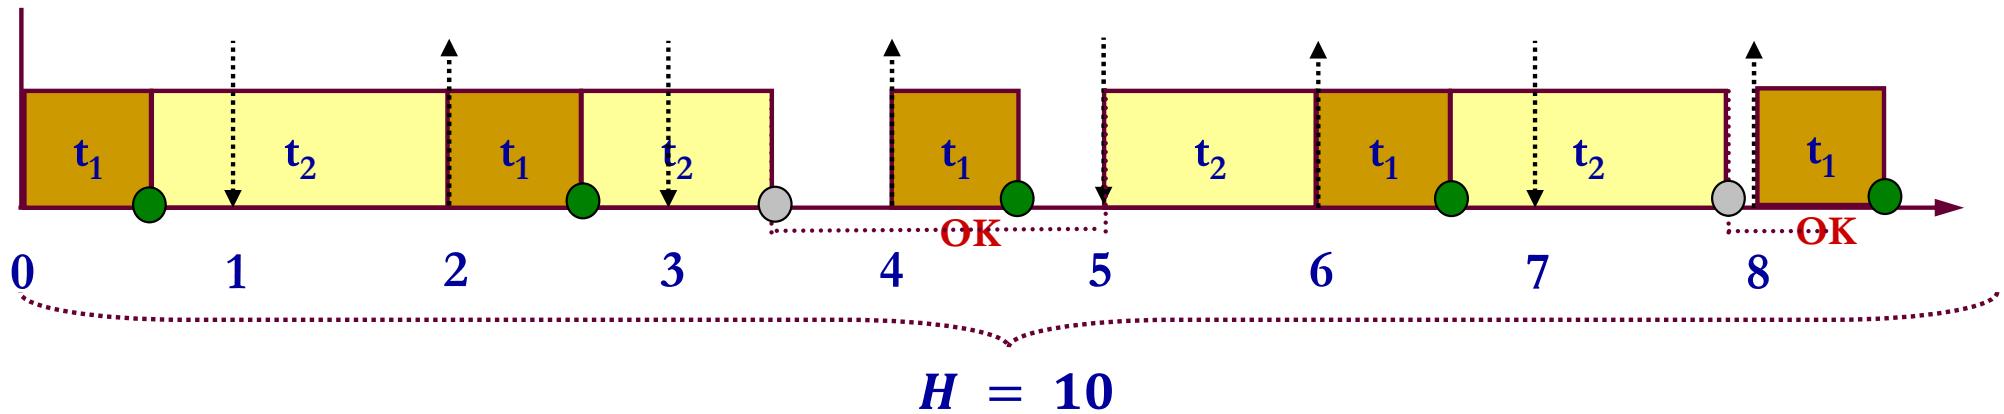
$$(\varphi_i, p_i, e_i, D_i)$$

$$T = \{\tau_1 = (0, 2, 0.6, 1), \tau_2 = (0, 5, 2, 3, 5)\}$$

$$\text{Density } \Delta(T) = \frac{e_1}{D_1} + \frac{e_2}{D_2} = 1.06 > 1$$

$$\text{Utilization } U(T) = \frac{e_1}{p_1} + \frac{e_2}{p_2} = 0.76 < 1$$

What happens to  $T$  under EDF?



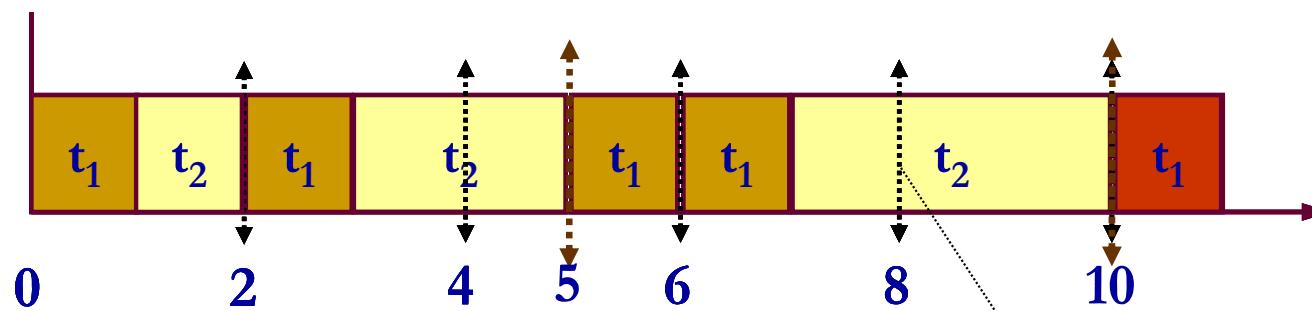
The exact utilization-based test tells us that  $T$  is feasible under EDF  
(We don't need to draw its timeline to tell that!)

# Overload situations – 5/5

$$(\varphi_i, p_i, e_i, D_i)$$

$$T = \{t_1 = (0, 2, 1, 2), t_2 = (0, 5, 3, 5)\} \Rightarrow U(t) = \frac{e_1}{p_1} + \frac{e_2}{p_2} = 1.1$$

T has *no* feasible schedule: what job suffers most under EDF?



$$T = \{t_1 = (0, 2, 0.8, 2), t_2 = (0, 5, 3.5, 5)\} \Rightarrow U(t) = \frac{e_1}{p_1} + \frac{e_2}{p_2} = 1.1$$

T has *no* feasible schedule: what job suffers most under EDF?

What about

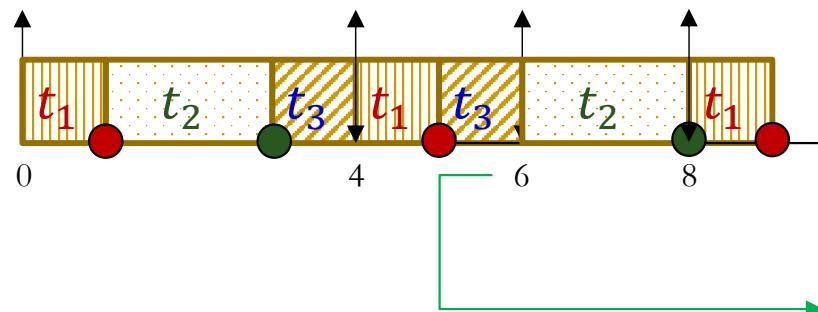
$$T = \{t_1 = (0, 2, 0.8, 2), t_2 = (0, 5, 4, 5)\} \text{ with } U(t) = \frac{e_1}{p_1} + \frac{e_2}{p_2} = 1.2 ?$$

# Preemption count – 1/2

$$(\varphi_i, p_i, e_i, D_i)$$

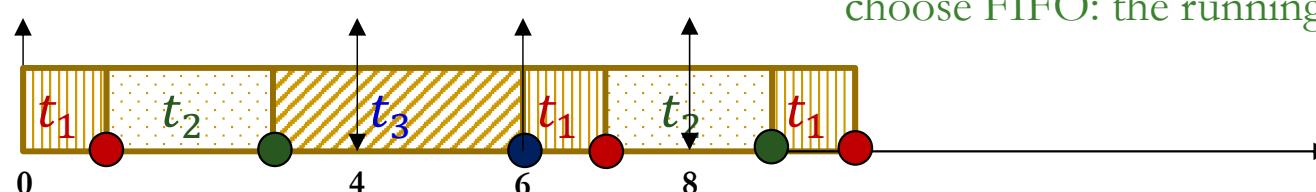
$$T = \{t_1 = (0, 4, 1, 4), t_2 = (0, 6, 2, 6), t_3 = (0, 8, 3, 8)\}, U = \frac{23}{24}, H = 24$$

With FPS and rate-monotonic priority assignment



With FPS, at time 4, with  
 $t_3$ 's absolute deadline = 8, priority = low  
 $t_1$ 's absolute deadline = 8, priority = high  
 $t_1$  preempts  $t_3$   
And, at time 6, with  
 $t_2$ 's absolute deadline = 12, priority = medium  
 $t_2$  preempts  $t_3$ , which misses its deadline  
(Remember: in case of equal priorities, EDF and FPS choose FIFO: the running job continues to run)

With EDF



EDF may incur *less* preemptions than FPS

# Preemption count – 2/2

Experiment	Run time	Mean preemptions FPS	Mean preemptions EDF	Min $\frac{P_{EDF}-P_{FPS}}{P_{FPS}}$	Max $\frac{P_{EDF}-P_{FPS}}{P_{FPS}}$
Fully-Harmonic	Hyperperiod	32,34	32,19	-0.5714	0.8571
Semi-Harmonic	Hyperperiod	4.265	4.255	-0.0282	0.1788
$1.0 < U < 1.0004$	Hyperperiod * U	23.385	41.171	-1.3866	-0.3089

Mean across task sets

# Back to FPS: critical instant – 1/2

- Feasibility and schedulability tests must consider the **worst case**, WC, for all tasks
  - The WC for task  $\tau_i$  occurs when the worst possible relation holds between its own release time and that of all higher-priority tasks
  - The actual case may differ depending on the admissible relation between  $D_i$  and  $p_i$
- The notion of **critical instant** – if one exists – captures the WC
  - The response time  $R_i$  for a job of task  $\tau_i$  with release time on the critical instant, is the longest possible value for  $\tau_i$

# Critical instant – 2/2

- **Theorem:** under FPS with  $D_i \leq p_i \forall i$ , the critical instant for task  $\tau_i$  occurs when the release time of *any* of its jobs is *in phase* with a job of every higher-priority task in the set
- We seek  $\max(\omega_{i,j})$  for all jobs  $\{j\}$  of task  $\tau_i$  for

$$\omega_{i,j} = e_i + \sum_{(k=1,..,i-1)} \left\lceil \frac{(\omega_{i,j} + \varphi_i - \varphi_k)}{p_k} \right\rceil e_k - \varphi_i$$

For task indices assigned in decreasing order of priority

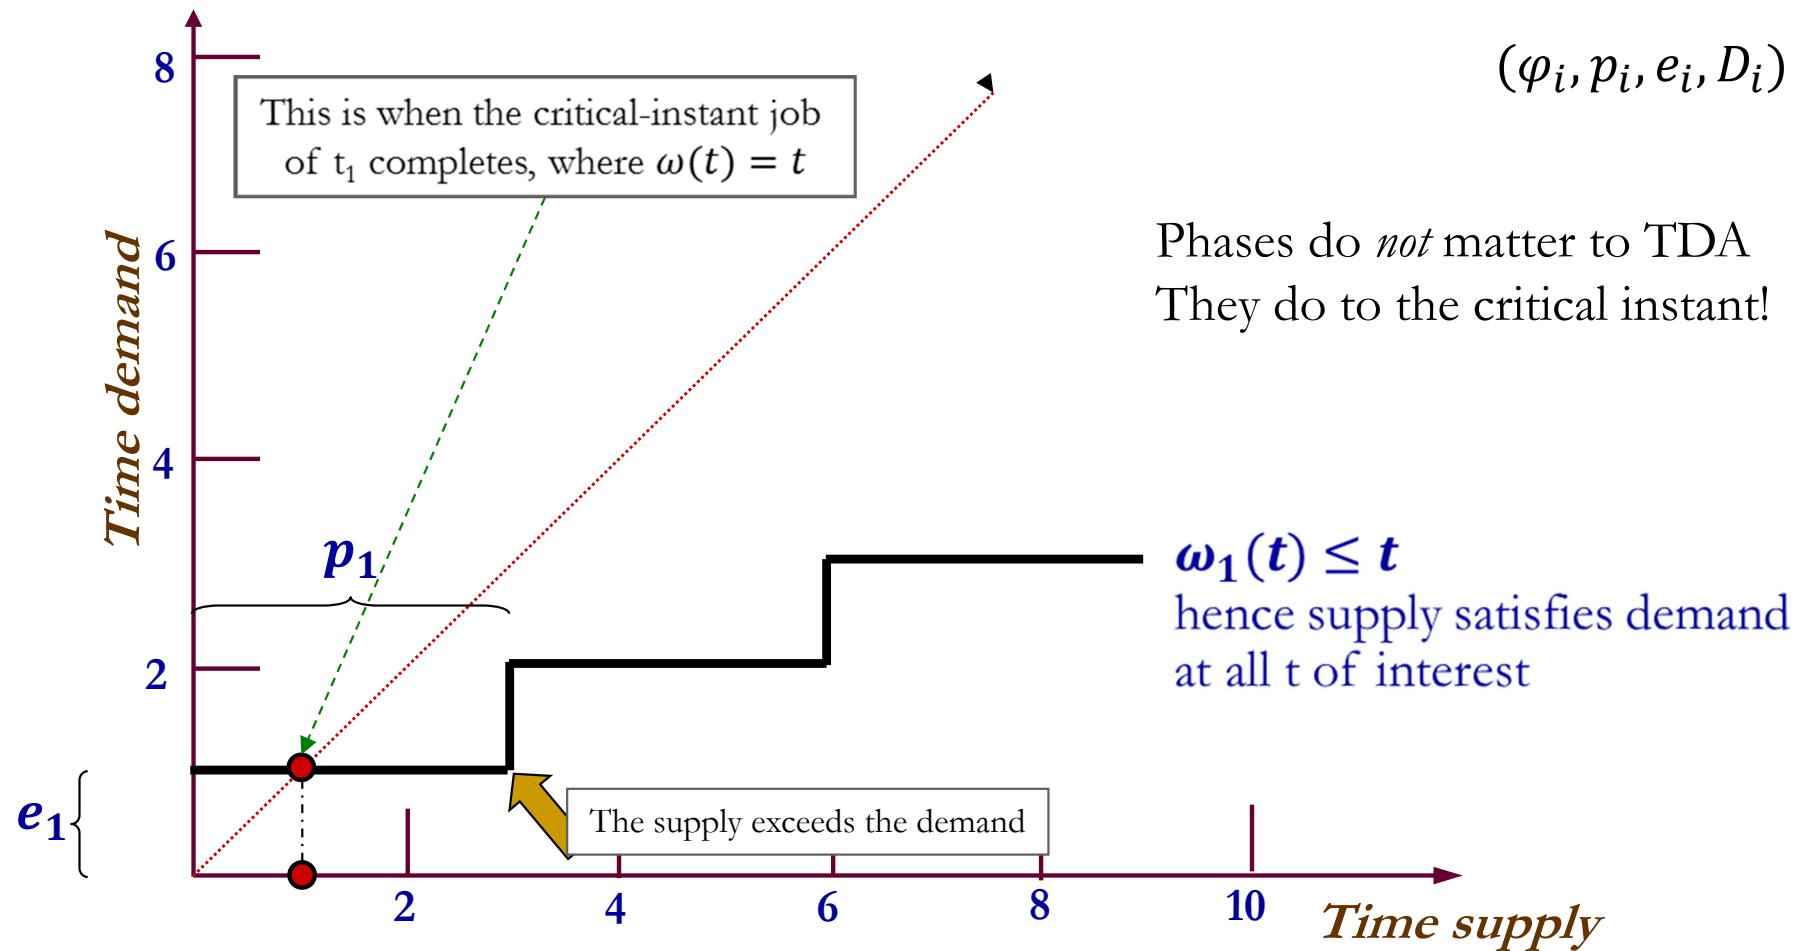
- The  $\sum$  component captures the **interference** that any job  $j$  of task  $\tau_i$  incurs from jobs of higher-priority tasks  $\{\tau_k\}$  between the release time of the first job of task  $\tau_k$  (with phase  $\varphi_k$ ) to the response time of job  $j$ , which occurs at  $\varphi_i + \omega_{i,j}$
- When  $\varphi$  is 0 for all jobs considered, all tasks are *in phase* and the equation captures the *absolute worst case* for task  $\tau_i$

# Time-demand analysis – 1/6

- ***Time Demand Analysis***, TDA, studies  $\omega$  as a function of time,  $\omega(t)$ 
  - As long as  $\omega(t) \leq t$  for some (selected)  $t$  for the job of interest, the supply satisfies the demand, hence the job can complete in time
- **Theorem** [Lehoczky, Sha, Ding: 1989]  
 $\omega(t) \leq t$  is an *exact feasibility test* for FPS
  - The obvious question is for which ‘ $t$ ’ to check
  - The method proposes to check at *all periods of all higher-priority tasks* until the deadline of the task under study

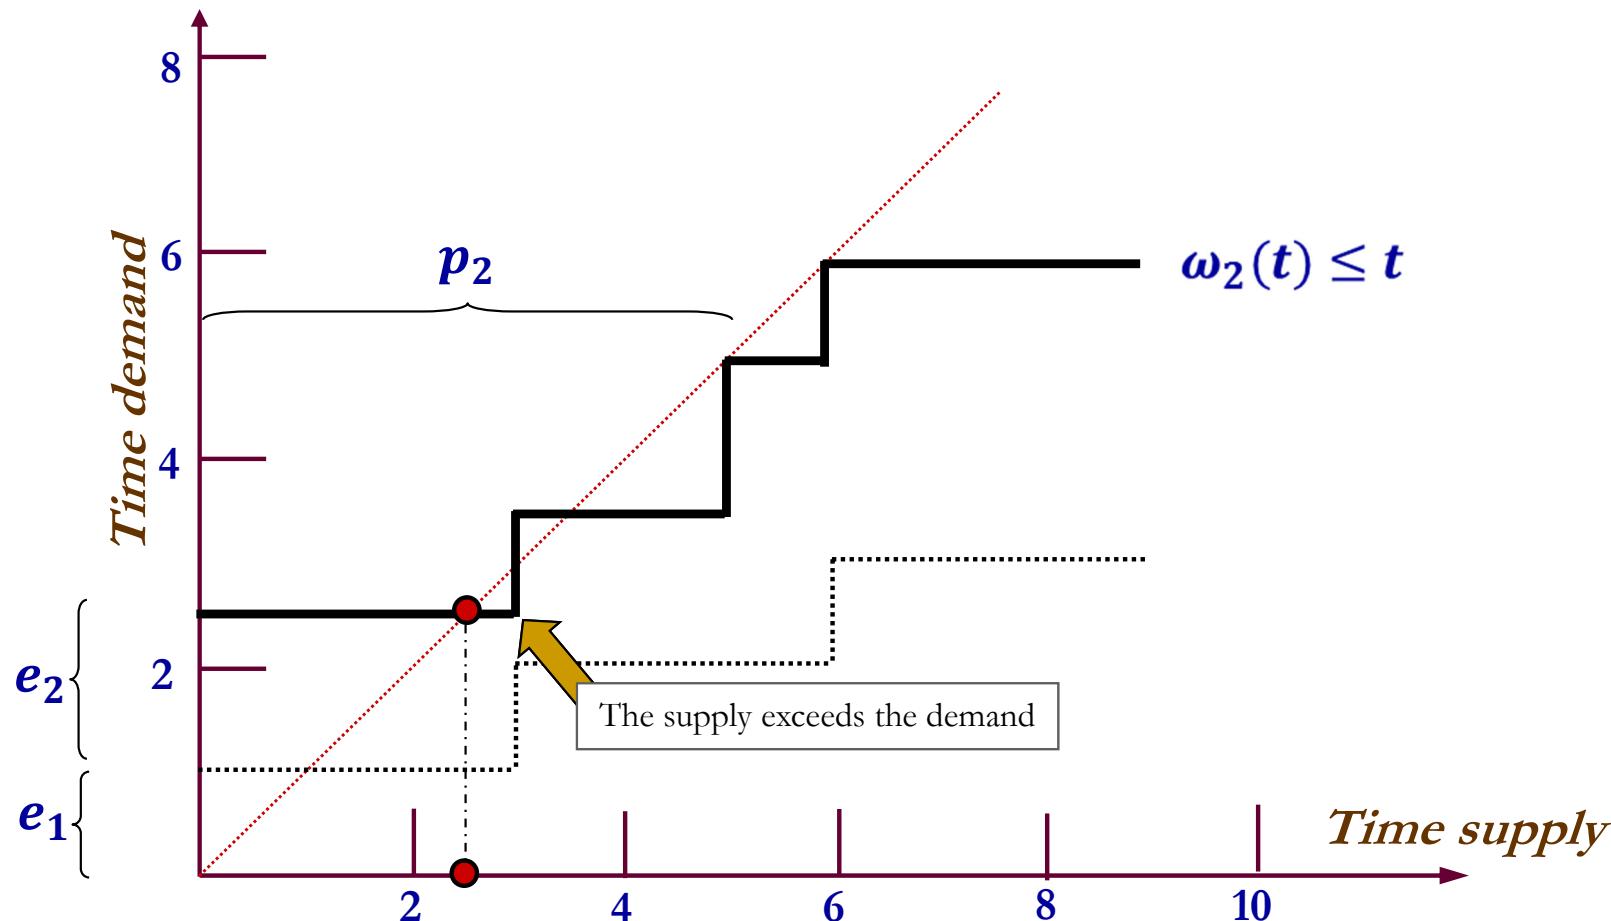
# Time demand analysis – 2/6

$$T = \{t_1 = (-, 3, 1, 3), t_2 = (-, 5, 1.5, 5), t_3 = (-, 7, 1.25, 7)\}, U = 0.82$$



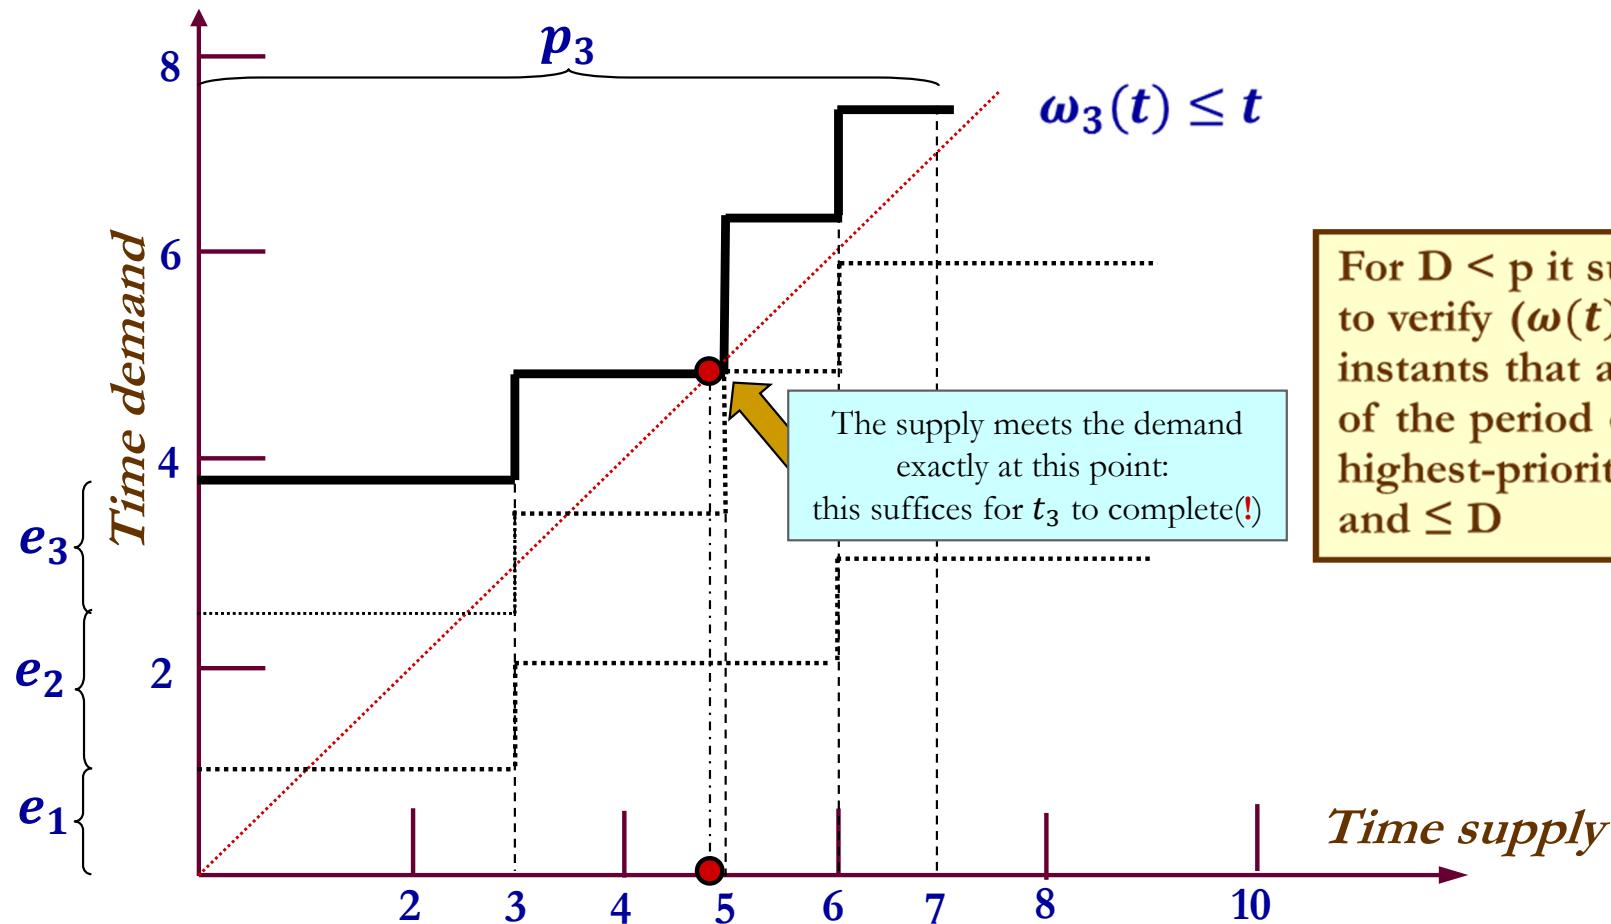
# Time demand analysis – 3/6

$$T = \{t_1 = (-, 3, 1, 3), t_2 = (-, 5, 1.5, 5), t_3 = (-, 7, 1.25, 7)\}, U = 0.82$$



# Time demand analysis – 4/6

$$T = \{t_1 = (-, 3, 1, 3), t_2 = (-, 5, 1.5, 5), t_3 = (-, 7, 1.25, 7)\}, U = 0.82$$



For  $D < p$  it suffices to verify  $(\omega(t) \leq t)$  at time instants that are multiple of the period of the highest-priority tasks and  $\leq D$

# Time demand analysis – 5/6

- We can use TDA to capture the *response time* of tasks and then use the critical instant notion to see that

The smallest value  $t$  that satisfies

$$t = e_i + \sum_{(k=1,..i-1)} \left\lceil \frac{t}{p_k} \right\rceil e_k$$

is the **worst-case response time** of task  $\tau_i$

- Solutions methods to calculate this value were independently proposed by
  - [Joseph, Pandia: 1986]
  - [Audsley, Burns, Richardson, Tindell, Wellings: 1993]

# Time demand analysis – 6/6

- **Theorem** [Lehoczky, Sha, Strosnider, Tokuda: 1991]  
When  $D > p$ , the first job of task  $\tau_i$  may *not* be the one that incurs the worst-case response time
- We must consider *all* jobs of task  $\tau_i$  within the so-called ***level-i busy period***, the  $(t_0, t)$  time interval within which the processor is busy executing jobs with priority  $\geq i$ , with release time in  $(t_0, t)$ , and response time falling within  $t$ 
  - The release time in  $(t_0, t)$  captures all backlog of interfering jobs
  - The response time of all jobs falling within  $t$  ensures that the busy period extends to their completion

# Example

$$\tau_1 = \{-, 70, 26, 70\}, \tau_2 = \{-, 100, 62, 120\}$$

$$(\varphi_i, p_i, e_i, D_i)$$

Let's look at the level-2 busy period

Ready queue:  $J_{1,1}, J_{2,1}$

**Time window 1 [0,70]**  
 Time left for  $J_{2,1}$ :  $70-26 = 44$   
 Still to execute:  $62-44 = 18$

Ready queue:  $J_{1,2}, J_{2,1}$

**Time window 2 [70,100]**  
 Time left for  $J_{2,1}$ :  $30-26 = 4$   
 Still to execute:  $18-4 = 14$   
 Release time of job  $J_{2,2}$

Ready queue:  $J_{2,1}, J_{2,2}$

**Time window 3 [100,140]**  
 Time left for  $J_{2,1}$  = 40  
 $J_{2,1}$  completes at: 114 ( $R = 114$ )  
 Time available for  $J_{2,2}$ :  $40-14 = 26$   
 Still to execute:  $62-26 = 36$

Ready queue:  $J_{2,2}, J_{2,3}$

**Time window 5 [200,210]**  
 Release time of job  $J_{2,3}$   
 $J_{2,2}$  completes at: 202 ( $R = 102$ )  
 Time available for  $J_{2,3}$ :  $10-2 = 8$   
 Still to execute:  $62-8 = 54$

Ready queue:  $J_{1,3}, J_{2,2}$

**Time window 4 [140,200]**  
 Time available for  $J_{2,2}$ :  $60-26 = 34$   
 Still to execute:  $36-34 = 2$

Ready queue:  $J_{1,4}, J_{2,3}$

**Time window 6 [210,280]**  
 Time available for  $J_{2,3}$ :  $70-26 = 44$   
 Still to execute:  $54-44 = 10$

Ready queue:  $J_{1,4}, J_{2,3}$

**Time window 7 [280,300)**  
 Time available for  $J_{2,3}$ :  $20-20 = 0$   
 Release time of job  $J_{2,4}$

Still in ready queue:  $J_{2,4}$   
 $\tau_2$ 's busy period  
 extends beyond  
 this point (!)

Ready queue:  $J_{1,5}, J_{2,3}, J_{2,4}$

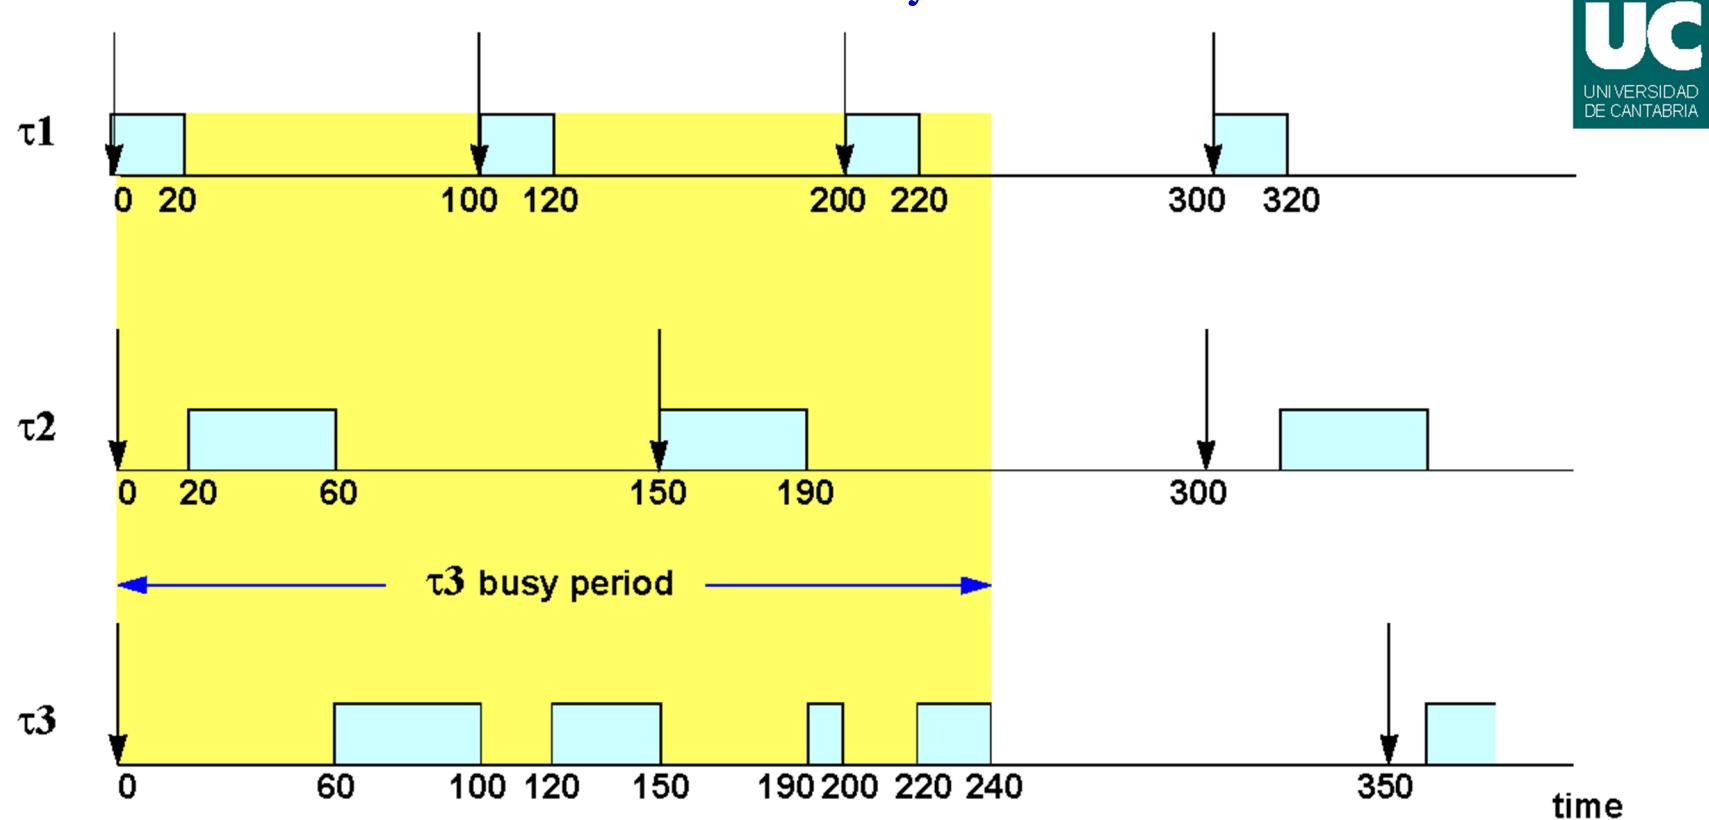
**Time window 8 [300,350)**  
 Time available for  $J_{2,3}$ :  $50-6 = 44$   
 $J_{2,3}$  completes at:  $300+6+10 = 316$  ( $R = 116$ )

$J_{2,1}$ 's response time is **not** worst-case!

# Level-i busy period

$$T_1 = \{-, 100, 20, 100\}, T_2 = \{-, 150, 40, 150\}, T_3 = \{-, 350, 100, 350\} \Rightarrow U = 0.75$$

The same definition of level-i busy period holds also for  $D \leq p$   
but its width is obviously shorter!



# Demand bound analysis (EDF)

- For  $\mathbf{df}$ , the EDF *demand function* and time  $t_i$ , an *exact* test for a task set  $T$  under EDF is:

$$\forall t_1, t_2: t_2 > t_i, \mathbf{df}(t_1, t_2) \leq t_2 - t_1$$

- For periodic tasks with no offsets and  $U \leq 1$ , it holds that:

$$\mathbf{df}(t_1, t_2) \leq \mathbf{df}(0, t_2 - t_1)$$

- The ***demand bound function*** helps generalize the test

$$\mathbf{dbf}(L) = \max_t(df(t, t + L)) = df(0, L), L > 0$$

- Theorem** [Baruah, Howell, Rosier: 1990] Exact test for EDF:

$$\forall L \in D(T), \mathbf{dbf}(L) \leq L, U < 1$$

- $D(T)$  is the set of deadlines for  $T$  in  $[0, L_m]$ ,  $L_m = \min(L_a, L_b)$ ,  $L_a = \max\{D_1, \dots, D_n, \frac{\sum_{i=1}^n (T_i - D_i) U_i}{1-U}\}$ ,  $L_b$  = first idle time in  $T$ 's busy period

# Summary

- Initial survey of scheduling approaches
- Important definitions and criteria
- Detail discussion and evaluation of main scheduling algorithms
- Initial considerations on feasibility analysis techniques

# Selected readings

- T. Baker, A. Shaw  
*The cyclic executive model and Ada*  
DOI: 10.1109/REAL.**1988**.51108
- C.L. Liu, J.W. Layland  
*Scheduling algorithms for multiprogramming in a hard-real-time environment*  
DOI: 10.1145/321738.321743 (**1973**)
- D. Perale, T. Vardanega  
*Removing bias from the judgment day: A Ravenscar-based toolbox for quantitative comparison of EDF-to-RM uniprocessor scheduling*  
DOI: 10.1016/j.sysarc.**2021**.102236